



Introduction to Infrastructure as Code




Deploying an Application

Provisioning




Topology of the
data center

Configuration



Setup and configure the
environment


Deployment



Deploy the application

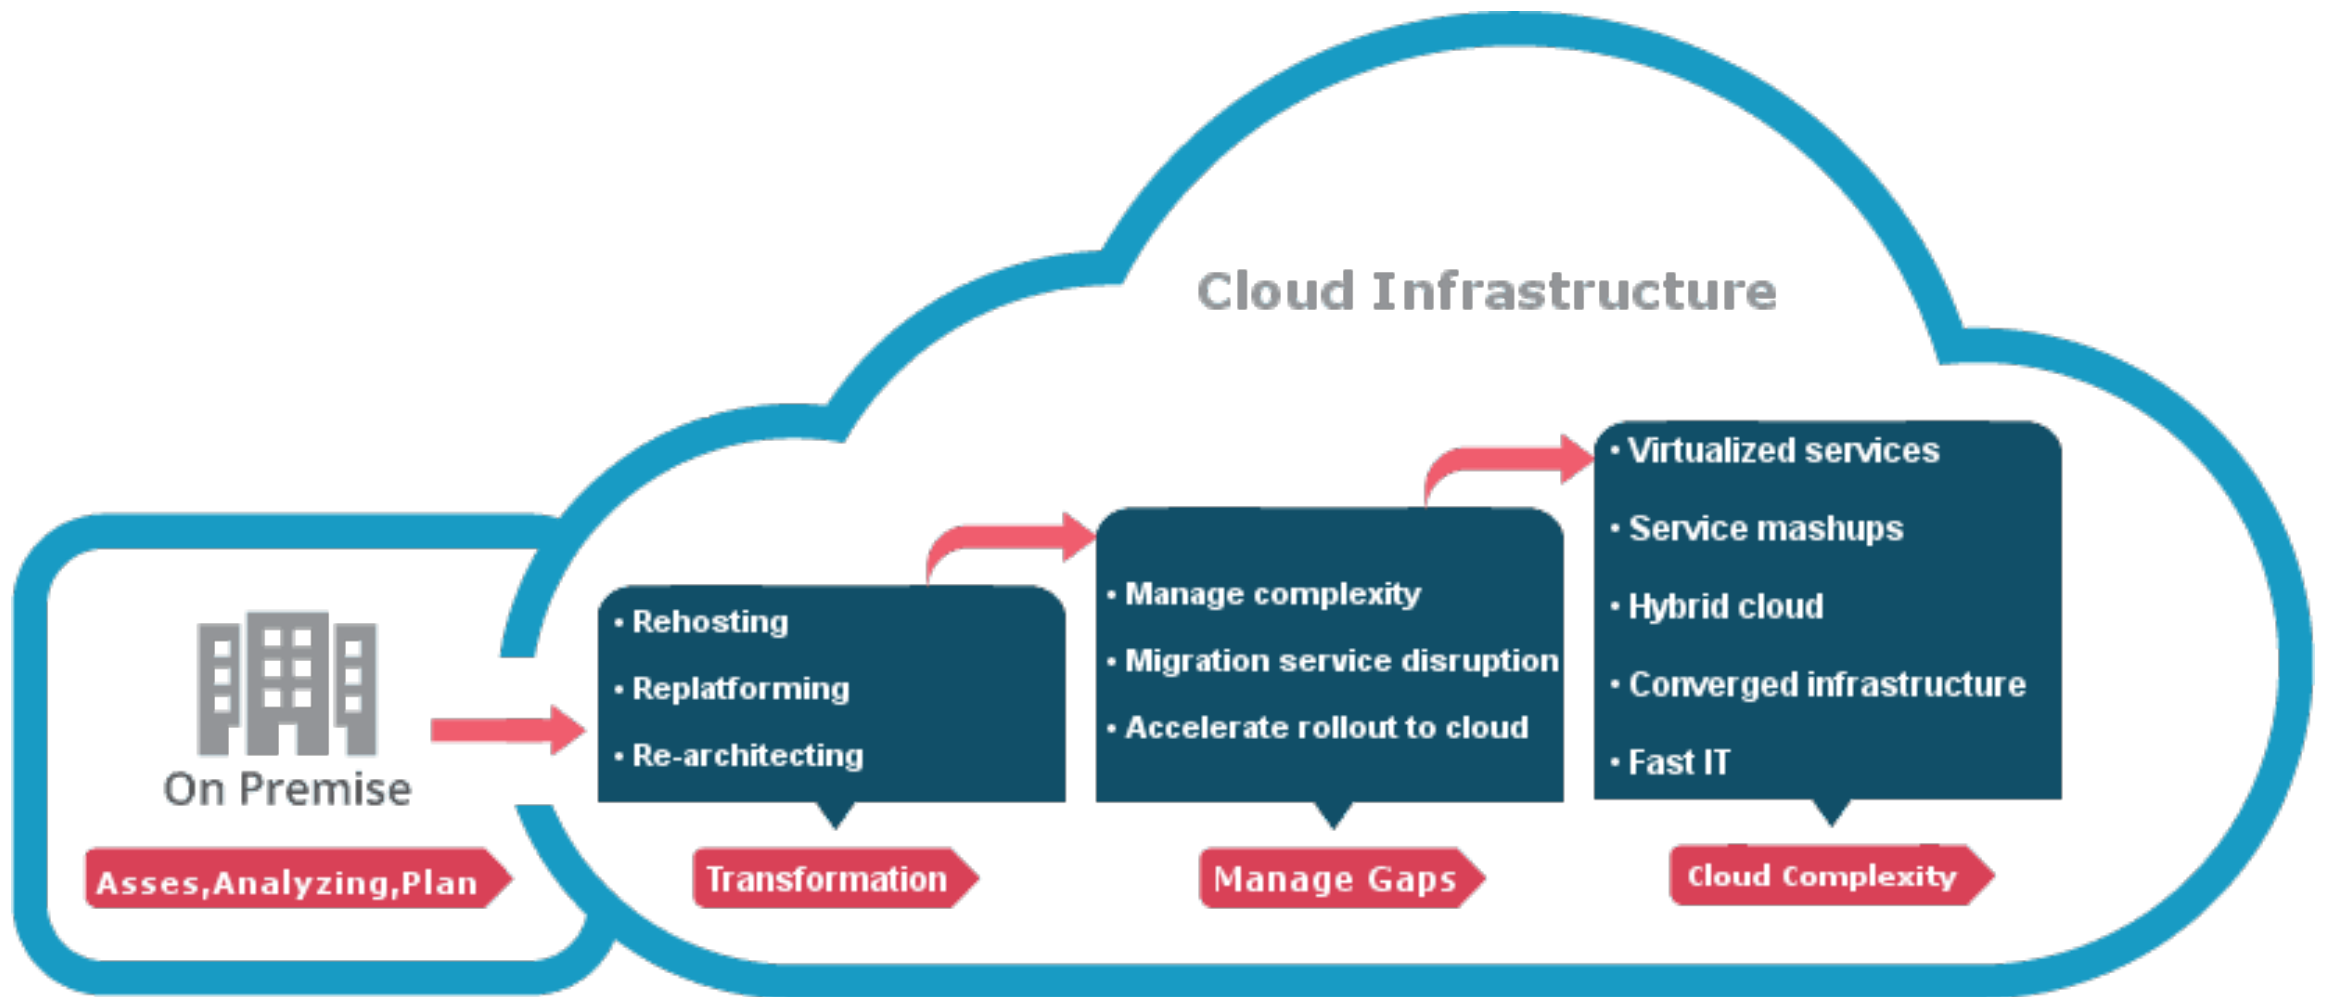


Provisioning Servers - The Old Way





Migrating to the Cloud





Typical Cloud Platform

Software as a Service

Managed applications and custom applications

Applications



Platform as a Service

Managed services eg. database, runtime, serverless, message queues, API gateway


Runtime and Services



Infrastructure as a Service

Traditional IT resources eg. servers, disk, subnets, cluster, backups

Infrastructure





Principles of Cloud Computing

- Assume system is unreliable
 - The hardware is unreliable because you have no visibility or control over them
 - Application must be reliable to compensate for the hardware's unreliability
- Create disposable things
 - Systems should be able to withstand dynamic changes to its constituent components
 - System components should be able to gracefully add, remove, stop, start, modify and migrate
- Make everything reproducible
 - Allow you to quickly recover from failures
 - Minimize statefulness especially in applications



Principles of Cloud Computing

- Minimize variation
 - Easier to automate the management if systems is largely the same eg. same OS, package versions, configuration, etc
 - Basis of building reproducible systems
- Ensure that you can repeat any action
 - Run and re perform any action will result in the same system/configuration
 - Capture list of actions in script, allow you to create systems with minimal variation



Principles of Cloud Computing

Immutable

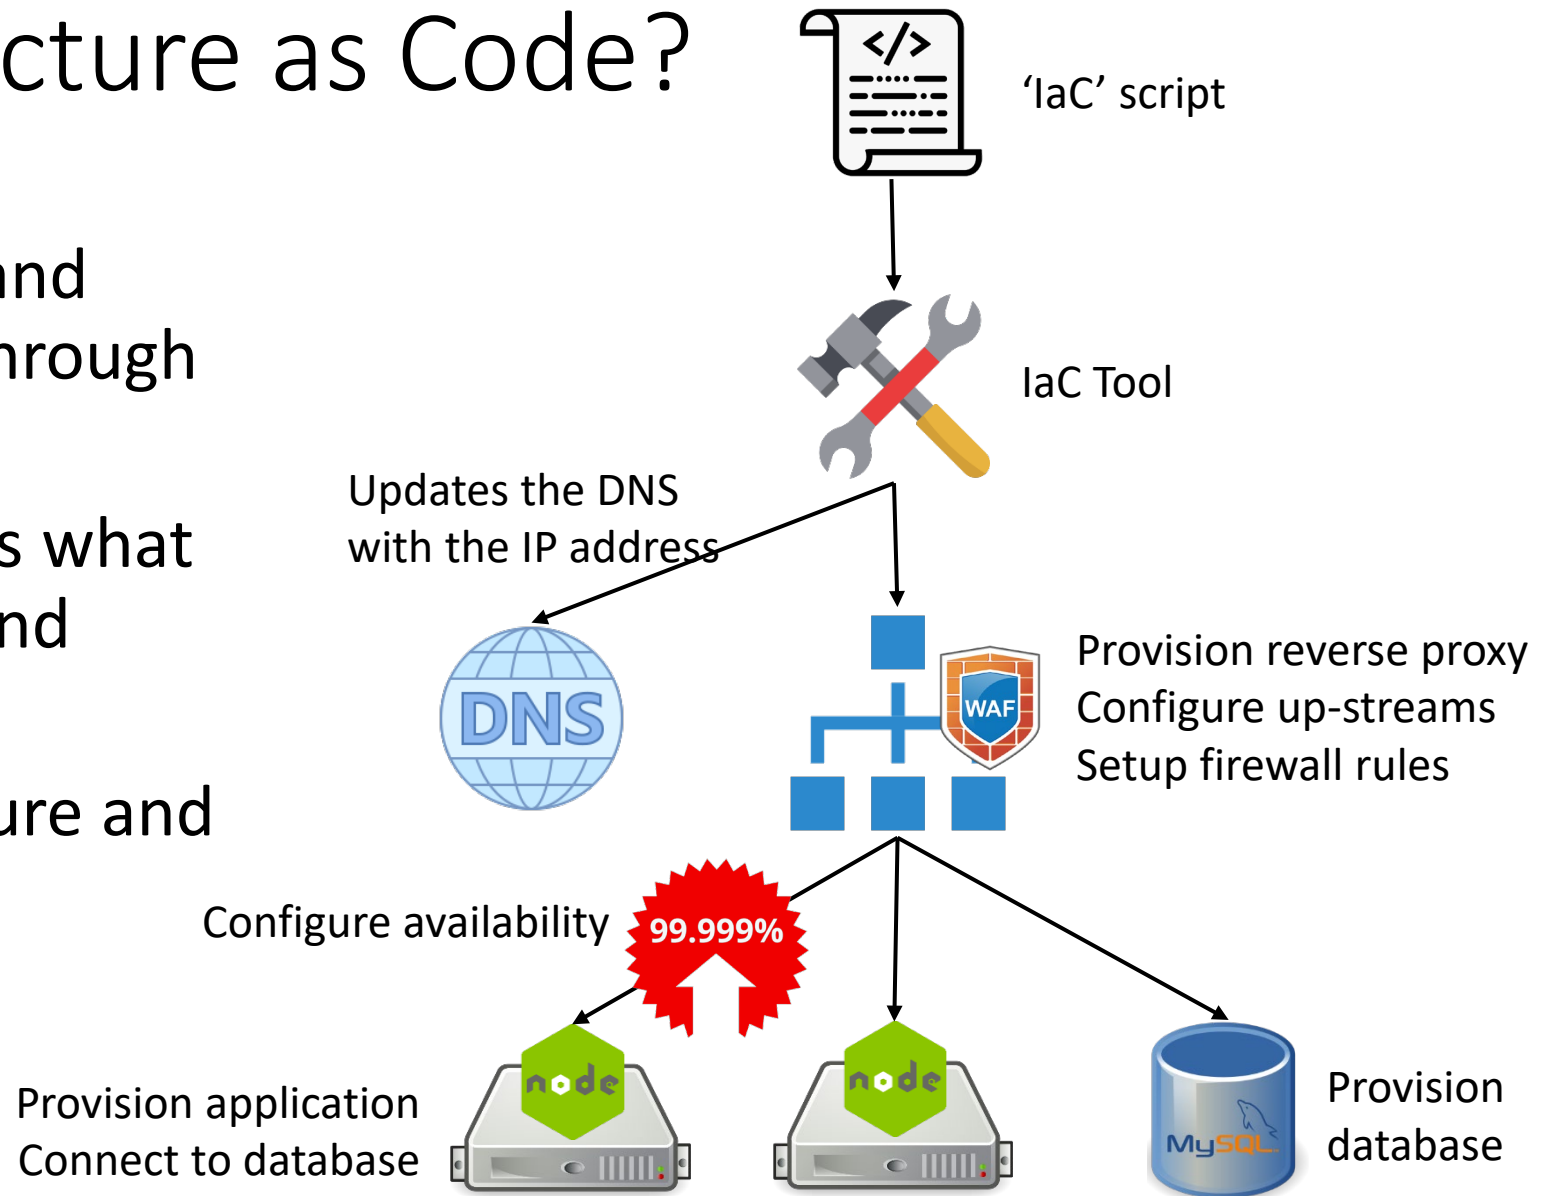
Ephemeral

Idempotent



What is Infrastructure as Code?

- Process of provisioning and managing IT resources through machine readable files
- Script/program describes what are the required setup and configurations
- Tool to provision, configure and install





Example of an IaC Script

```
resource "digitalocean_ssh_key" "web_ssh_key" {  
  name = "web_ssh_key"  
  public_key = file("./mykey.pub")  
}
```

Provision a server with a SSH key
Install Nginx

```
resource "digitalocean_droplet" "webapp_droplet" {  
  image = "ubuntu-20-40-x64"  
  name = "webapp_droplet"  
  region = "sgp1"  
  size = "s-1vcpu-1gb"  
  ssh_keys = [ digitalocean_ssh_key.web_ssh_key.fingerprint ]  
  provisioner "remote-exec" {  
    inline = [  
      "apt update",  
      "apt upgrade -y",  
      "apt install nginx -y"  
    ]  
  }  
}
```



Why is this Possible?

- IT resources are mostly virtualized
 - Through cloud providers, Kubernetes/Docker
- Result of virtualization, resources can be provisioned in minutes via browser or command line

```
doctl compute droplet create \  
  --image ubuntu-20-04-x85 --size s-1vcpu-1gb \  
  --region sgp1 myserver
```

- Expose endpoint to support tools and applications
- Developers can leverage these API
 - Through official SDKs
 - Calling the endpoints directly, typically HTTP



Benefit of IaC - Immutable Infrastructure

- Infrastructure is codified in scripts
 - Will deploy the same exact setup every time the script is executed
 - Captures know-how
- Mitigate configuration drift
 - “Works on my machine”
 - Consistent updates, server configurations the same across all servers
 - No snowflakes or fragile infrastructure
- Parity between development, testing and deploy
 - Fewer issues when moving from development to testing to deployment
 - More confident when application goes into production because the testing environment is the same as the production
- Promotes infrastructure reuse
 - Modularize and parameterize the infrastructure eg. what goes into a basic server
 - Consistently use the same basic building blocks, increase security and availability



Benefits of IaC - Programmers Paradigm



- Tools to help with script development
 - Syntax checking, code completion, linting
- Abstraction and reuse
 - Package commonly used infrastructure as 'components' to be shared
- Testing
 - Unit, integration, smoke, end to end
- Tracked and versioned controlled
 - Reviewed via pull request



Result of Adopting IaC

- Decrease the time for implementing, testing and deploying application/infrastructure changes into production
 - Easier to setup consistent environment for testing
 - Bug fixes, new feature releases can easily be tested
- Increase deployment frequency
 - Environment can be created easily, can lead to more frequent releases for new features
 - Align with agile practice - small frequent updates



Result of Adopting IaC

- Reduce the number of failures when an change is applied to a production application
 - Parity of environment in development, testing, staging and production
 - Reduce the probability of misconfiguration resulting in failure
- Reduce the time it takes to restore an application
 - Entire application's configuration and setup is describe and capture in scripts
 - Application can be recreated by running the scripts
 - Note: Need to factor data restoration time



Infrastructure as Code - Mindset Changes

Manual Way	Automation Way
Cost of changes are high	Cost of changes are low
Changes represent failure, therefore must be managed and controlled	Changes represent learning and improvement
Reduce opportunity to fail	Maximize speed of improvement
Application features are delivered in large batches of changes, test end-to-end	Deliver small changes, test continuously
Applications have long release cycle	Applications have short release cycle
Monolithic architecture - larger, less moving parts	Microservices architecture - smaller, more moving parts
Physical or GUI driven changes	Code as configuration



IaC Tools



Ansible: installing and managing the servers



Packer: write an image



Infrastructure Templating

Manage infrastructure

Install applications and
one time configurations

Deploy configuration and changes
post install




IaC – Declarative vs Imperative

- Two ways of writing IaC scripts
 - Declarative – define the final state of the infrastructure, tool will determine how best to arrive at the final state from the present state
 - Imperative – specify the steps to arrive at the final state, the tool will follow the steps to provision the required resources
- Scripts may be written in
 - Domain specific language (DSL)
 - Eg. Terraform, Ansible, Kubernetes
 - General purpose programming language
 - Eg. Pulumi, AWS CDK in Golang, Java, JavaScript, Python, etc

Declarative




Define the end state
Tool will reconcile the difference



Final state

3 servers
1 disk



Specify the steps to reach the end state
Tool will follow the steps

Imperative



Example: IaC with Programming Language

```
import * as awsx from "@pulumi/awsx";
import * as pulumi from "@pulumi/pulumi";

// Create an elastic network listener to listen for requests
// and route them to the container.
const listener = new awsx.elasticloadbalancingv2.NetworkListener("nginx", { port: 80 });

// Define the service to run. Pass the listener to hook up the network load balancer
// to the containers the service will launch.
const service = new awsx.ecs.FargateService("nginx", {
  desiredCount: 3,
  taskDefinitionArgs: {
    containers: {
      nginx: {
        image: awsx.ecs.Image.fromPath("nginx", "./app"), memory: 512,
        portMappings: [listener] }
    }
  }
});
export let frontendURL = pulumi.interpolate `http://${listener.endpoint.hostname}/`;
```




Idempotent

- Property of an operation where they can be applied many times without changing the beyond the initial application
- IaC should be idempotent
 - Will not deploy more resources when the final state has been reached
 - Some imperative operations may not be idempotent
- Tools required to maintain the state of the resources under its management
 - Compare the actual versus the desired
 - Can be a potential source of contention especially when working in a team




Maintaining State

- Need some way to understand the state of the deployment for idempotency
- State reconciliation is either done via push or pull
- States can be shared across a team
 - Some tools support state locking
 - Prevent multiple updates that can corrupt the state






What is DevOps?



- DevOps is a combination of culture, practices and tools
- Aim is to shorten the systems development lifecycle
 - By streamlining software building, testing and release
- Benefit is the improve and evolve applications at a faster pace
- Automation is they key to the fast feedback/delivery loop
 - Key tool is CICD



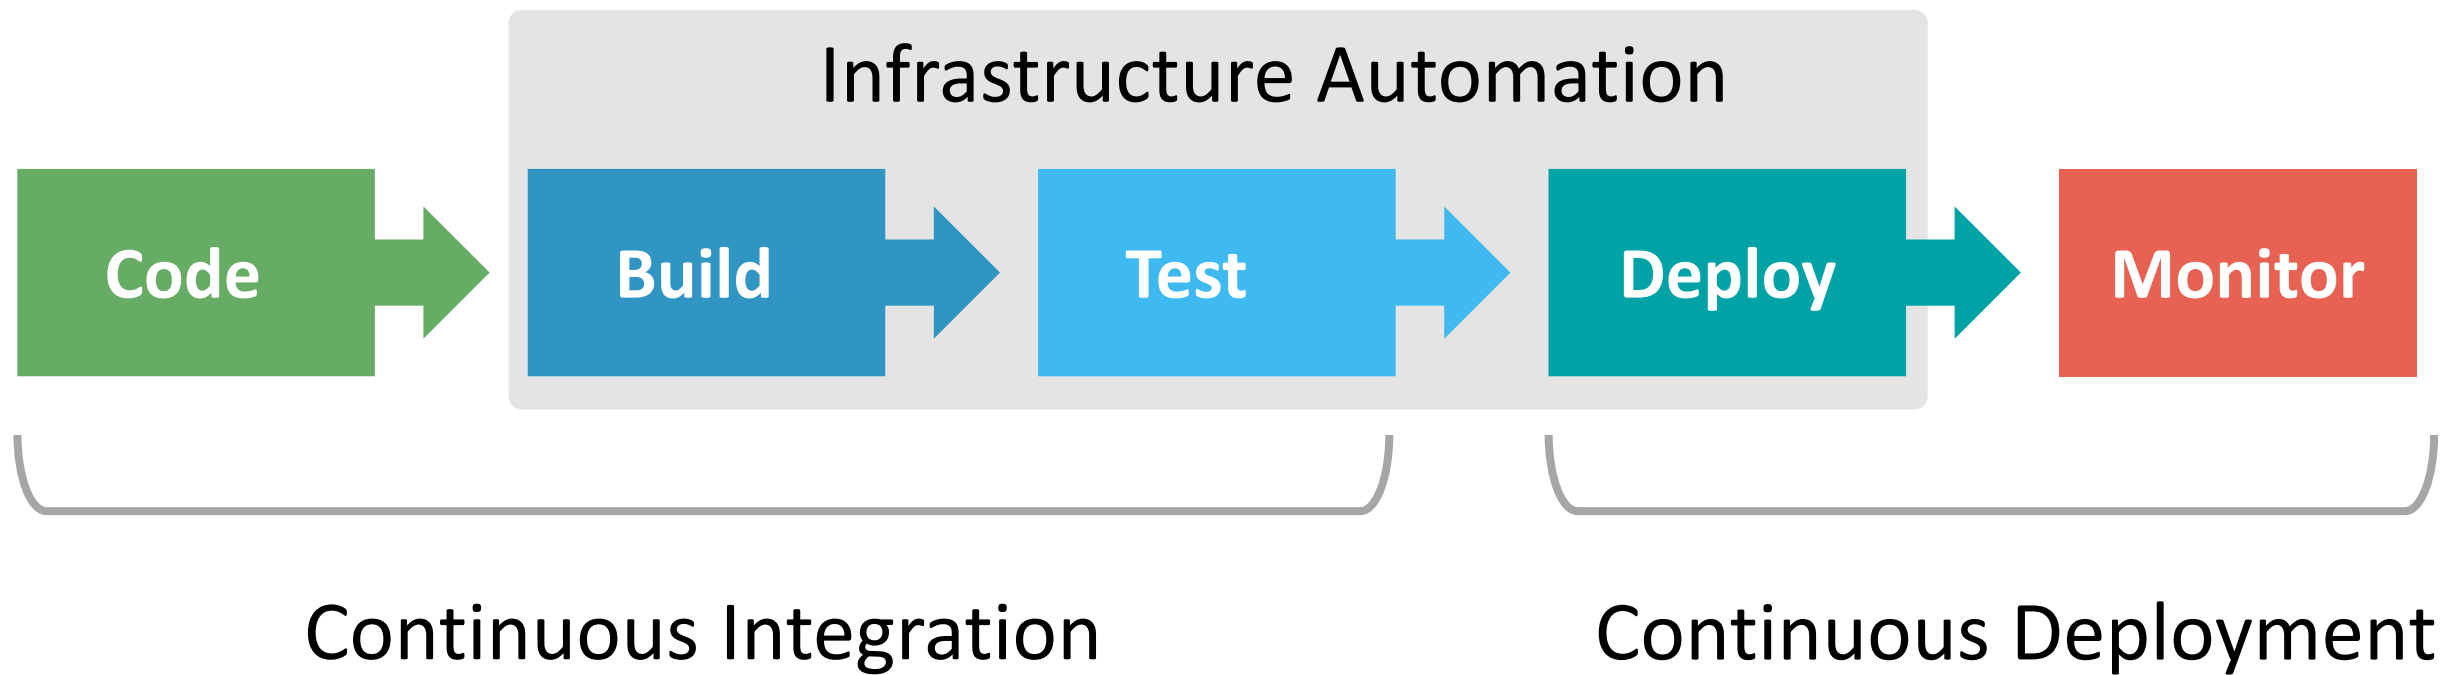
What is CI/CD?



- Refers to the practice of continually building and delivering an application
 - 2 phase process
- Continuous Integration where modules that make up an application are continuously tested and integrated
- Continuous Delivery is the practice of ensuring that the code is always at a deployable state
- Leverage tools to automate the integration and deployment



Infrastructure as Code in CI/CD Pipeline





Appendix



Testing

- Software testing focus on the application
 - Unit, stress, smoke, end-to-end
 - Testing known conditions
 - Test if the application conforms the requirements and data flow
- What about unknown or unforeseen errors in infrastructure, application, processes
 - Eg misconfigured database cluster
 - Eg. the application itself is down due to hardware failure
 - Eg. data corruption
- 'Regular' testing do not test these scenarios



What is Chaos Engineering?

- Chaos engineering is the discipline of experimenting on a system in order to build confidence in the system's capability to withstand turbulent conditions in production.
 - From <https://principlesofchaos.org>
- Idea is to break your system on purpose at random time during production
 - Inject fault into the system like vaccine, inject into body to build immunity
 - Observe how the infrastructure, application, process and people react and respond
- Reveal readiness and build confidence in handling crisis



Build Resilient System

Software

- Memory leaks
- License
- Certificate expiration
- Incompatible version

Infrastructure

- Redundancy
- Self healing
- Isolation
- Infrastructure as Code

Application

- Request timeouts
- Data corruption
- Rate limiting
- Exception handling

Operations

- Logging
- Monitoring and observability
- Incident response
- Metrics




DYI Simple Fault Injection

- Stop the database
 - `sudo systemctl stop mysql`
- DDoS your application
 - `siege -c 100 https://myapp.com`
- Max out your system resources
 - `stress-ng --cpu 8 --vm 4`
- Randomly drop packets on network interface
 - `tc qdisc add dev eth0 root netem loss 20%`
- Block access to DNS
 - `iptables -A INPUT -p tcp -m tcp --dport 53 -j DROP`



Chaos Engineering Toolkit

- Chaos Monkey is a set of tools to implement chaos engineering
 - Developed by Netflix
 - <https://github.com/netflix/chaosmonkey>
- Set of agents that performs the following
 - Shuts down services randomly
 - Slows down performance
 - Can be scheduled
- Deployable on Kubernetes
 - Other options: Kubernetes 'native' - Chaos Mesh





Infrastructure Elements as Code

- Server image definition
 - Operating system
- Software packages to be applied to server
 - What additional packages to be installed on a generic server
- Application artefacts eg. binaries, sources, configurations
- Infrastructure stack
 - Collection of servers, networks, API gateways
- Operation services configurations
 - Firewall rules, API gateway healthchecks,
- CI/CD pipeline definitions
 - Include testing and validation



Idempotent

- Property of an operation where they can be applied many times without changing the beyond the initial application
- Most IaC are idempotent
 - Will not deploy more resources when the final state has been reached
 - Some imperative operations may not be idempotent
- Tools required to maintain the state of the resources under its management
 - Compare the actual versus the desired
 - Can be a potential source of contention especially when working in a team