

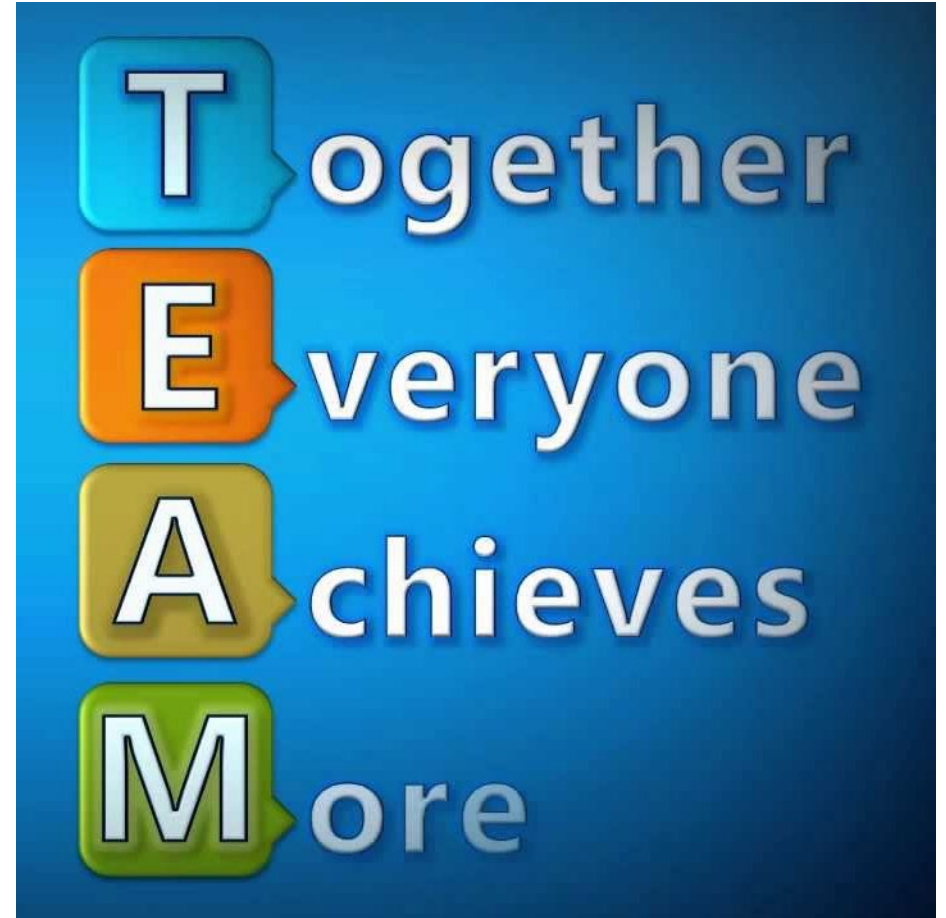


# Working in a Team



# Working in a Team

- Not always easy or transparent
  - Whatever the group is doing is opaque to every other group
  - Result in unstable or inconsistent environment
- Fear of adopting automation tools
  - Especially tools that apply changes to existing systems
  - Not sure what might break
- Making changes easily and safely (IaC context)
  - Conventional wisdom is to sacrifice one for other





# Terraform Modules

- Reusable infrastructure/architectural pieces
  - Any Terraform project is a module
- Module has
  - A set of parameters - eg the number of server instance to provision
  - A set of outputs - eg IP address of all the servers
  - Internal states - local variables, defined and used by the module
- Modules can be shared by publishing to public repository
- Terraform files inside a sub-directory becomes a module
- Modules require their own providers
  - Even if the root module is using the same provider

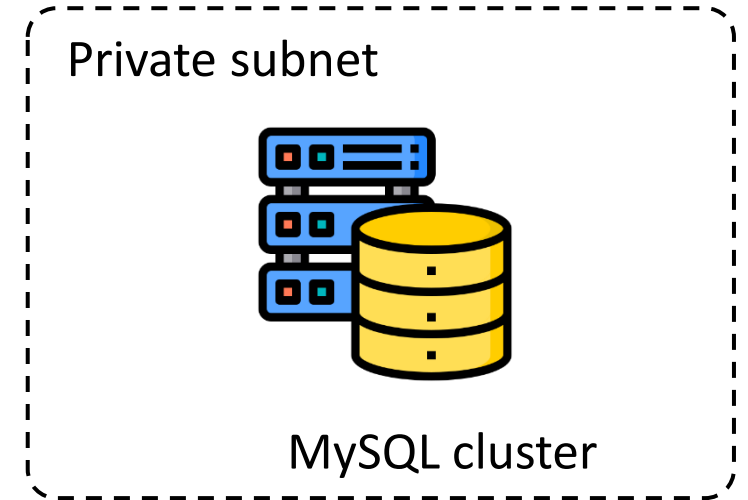


# Example - Module

resources.tf

```
resource digitalocean_vpc vpc {  
  name = "${var.name}_vpc"  
  region = var.region  
}
```

```
resource digitalocean_database_cluster db {  
  name = "${var.name}_db"  
  region = var.region  
  size = var.db_size  
  engine = "mysql"  
  version = "8"  
  node_count = var.node_count  
  private_network_uuid = resource.digitalocean_vpc.vpc.id  
}
```



Module will also need a provider configuration. Not shown



# Example - Module

## inputs.tf

```
variable region {
  type = string
  description = "Region to ..."
}

variable name {
  type = string
  description = "VPC name"
}

variable db_size {
  type = string
  description = "Database instance"
  default = "db-s-1vpc-2gb"
}

...
```

## outputs.tf

```
output vpc_id {
  value = digitalocean_vpc.vpc.id
  description = "VPC id"
}

output db_connection {
  value = digitalocean_database_cluster.db.uri
  description = "MySQL connection string"
}

output db_user {
  value = digitalocean_database_cluster.db.user
  description = "MySQL user"
}

output db_user {
  value = digitalocean_database_cluster.db.password
  description = "MySQL password"
  sensitive = true
}
```



# Example - Module

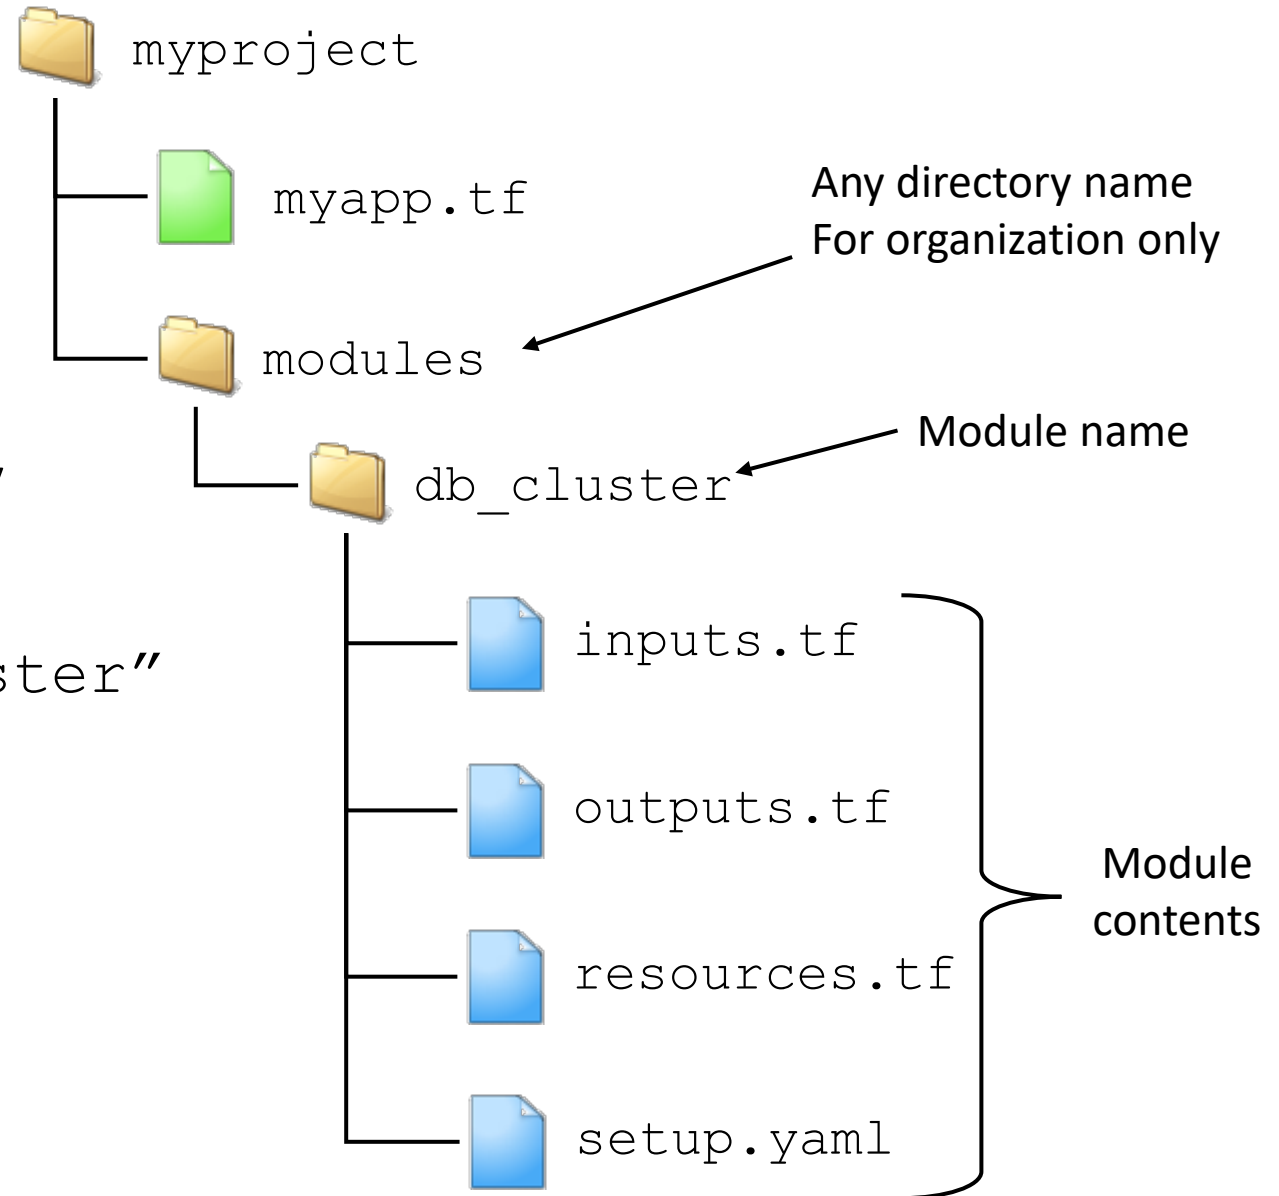
Module instance name  
myapp.tf

```
module myapp {  
  source = "../modules/db_cluster"
```

Location of the module directory

```
  name = "myapp"  
  region = "sgp1"  
  node_count = 2  
}
```

Module configurations  
input variables



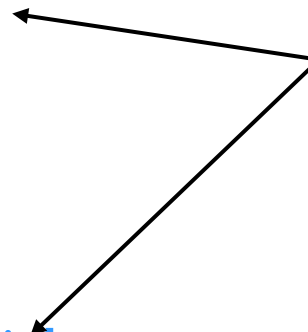


# Example - Module

```
resource digitalocean_droplet mydroplet {  
  ...  
  vpc_uuid = module.myapp.vpc_id  
}
```

```
output db_connection {  
  value = module.myapp.db_connection  
}
```

Accessing the  
outputs of a module





# Local Variables

- Local variables are used within the module
  - Not accessible by any resources outside of the module
- Local variables are access internally within the module with the `local` prefix

locals.tf

```
locals {  
    db_username = "fred"  
}
```

resources.tf

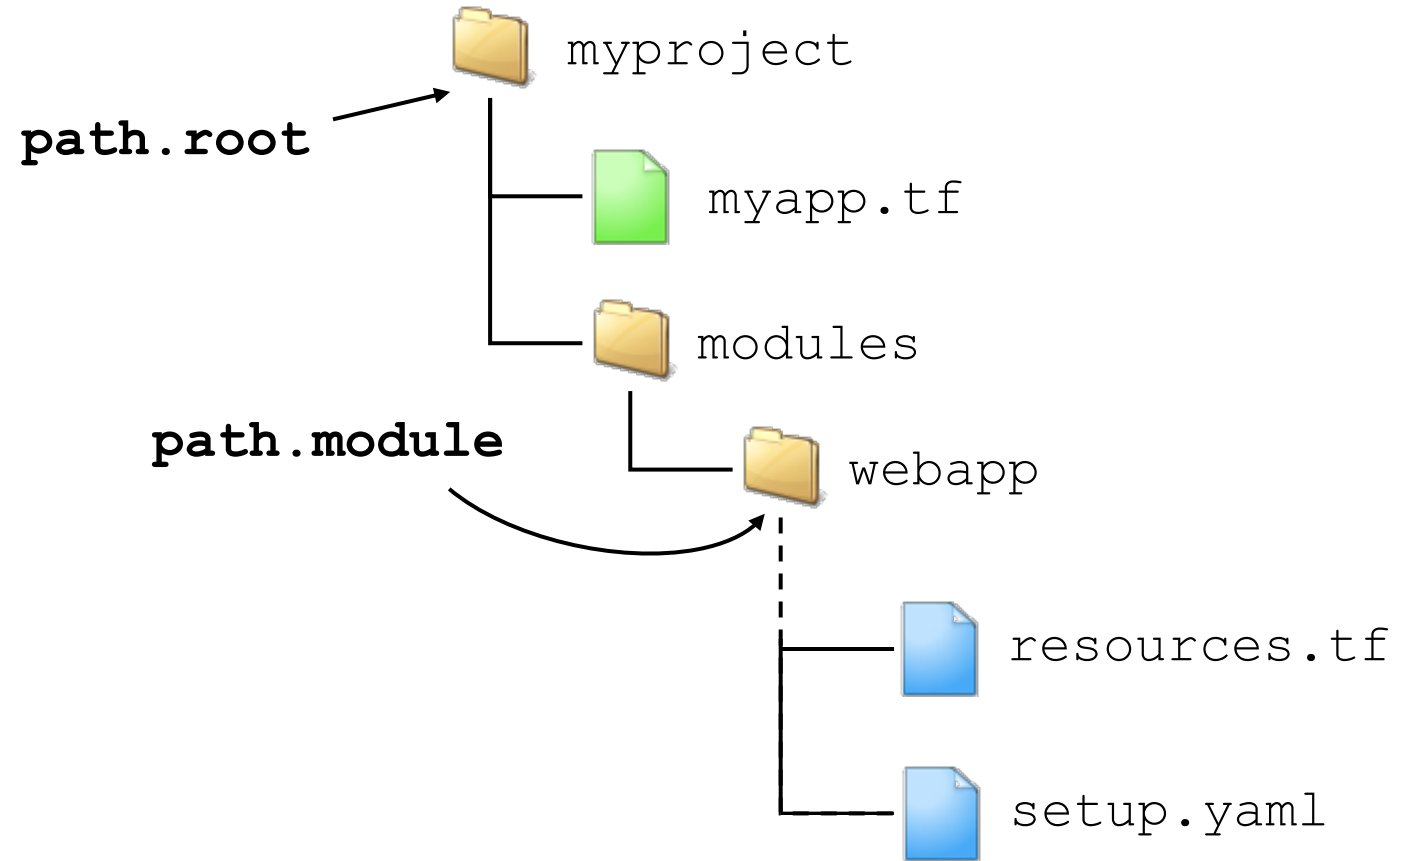
```
resource digitalocean_database_user dbuser {  
    cluster_id = digitalocean_database_cluster.db.id  
    name = local.db_username  
}
```





# File Path

- Modules provides the following method to transparently reference files relative to its location
  - `path.module` - the path to the module root
  - `path.root` - the path of the root module viz. the module that is using



```
resources.tf
resource digitalocean_droplet droplet {
  name = var.droplet_name
  image = var.droplet_size
  region = var.droplet_region
  size = var.droplet_size
  user_data = file("${path.module}/setup.yaml")
}
```



# State File

- State file records the resources provisioned by Terraform
  - Created when you run `terraform apply`
- Maps the Terraform resource with the actual resource
  - Consult the state file whenever you run `terraform plan`
  - Reconcile what is defined in the `.tf` files vs the real world
  - Only provisioned those missing resources when apply is executed



# State File

- Shared location of the state file
  - State files are created locally, so not accessible to another member in the group
  - Other team members can provision duplicate resource without access to a common shared state file
  - Simple solution is to save state file in a shared network accessible storage
- Corrupting state file with multiple simultaneous updates
  - Race condition, may cause data loss or corruption of the state file
- Single state files for multiple different environment
  - A state file only store the state of one environment eg. testing
  - Members in the same team cannot update the resources, may accidentally cause outage or failure in the test



# Backends

- Pluggable state file storage
  - Can be configured to create, read and update state file from different backend storages
- Default backend is local
  - State file is created in the current directory, viz. `path.root`
- Other backends include S3, Postgres, Kubernetes, HTTP (WebDAV), EtcdV3,
  - Some backend supports locking
  - See <https://www.terraform.io/docs/language/settings/backends/index.html>
- Can only configure 1 backend per project



# Example - local

```
terraform {  
  required_version = ">= 0.15.0"  
  required_providers = {  
    digitalocean = { ... }  
  }  
}
```

```
  backend local {  
    path = "path/to/file/terraform.tfstate"  
  }  
}
```

Specify the backend to use.  
local is the default

Change the default  
location of the state file

The file can be stored in local hard disk or a  
shared directory like NFS or Gdrive  
local does not provide file locking



# S3 Backend

- Backend supports all AWS backend features including versioning
  - Best practice to enable versioning for S3, for recovering old state files in the event of an accident
- S3 bucket needs to have actions
  - `s3:ListBucket`, `s3:GetObject`, `s3:PutObject`
  - Either for a specific principal for the backend or all
- Supports locking by configuring a DynamoDB table name
- Table needs to have the following actions
  - `dynamodb:GetItem`, `dynamodb:PutItem`, `dynamodb.DeleteItem`



# Example - S3 Compatible Storage

```
terraform {  
  ...  
  
  backend "s3" {  
    skip_credentials_validation = true  
    skip_metadata_api_check    = true  
    skip_region_validation     = true  
    skip_requesting_account_id = true  
    skip_s3_checksum           = true  
    endpoints {  
      s3 =  
"https://sgp1.digitaloceanspaces.com"  
    }  
    region = "sgp1"  
    bucket = "mybucket"  
    key    = "states/terraform.tfstate"  
  }  
}
```

Skip all validation for S3  
compatible storage



- Standard S3 parameters
- Access key and secret access key for S3 bucket
  - Can be configured as arguments in the backend (not recommended)
  - As environment variable  
AWS\_ACCESS\_KEY\_ID and  
AWS\_SECRET\_ACCESS\_KEY
  - Read from  
~/.aws/credentials
- Need to disable validation if using S3 compatible storage
  - Eg. DigitalOcean, Cloudflare



# Reconfigure Backend

- Need to reconfigure the backend if it changes

```
terraform init -reconfigure
```



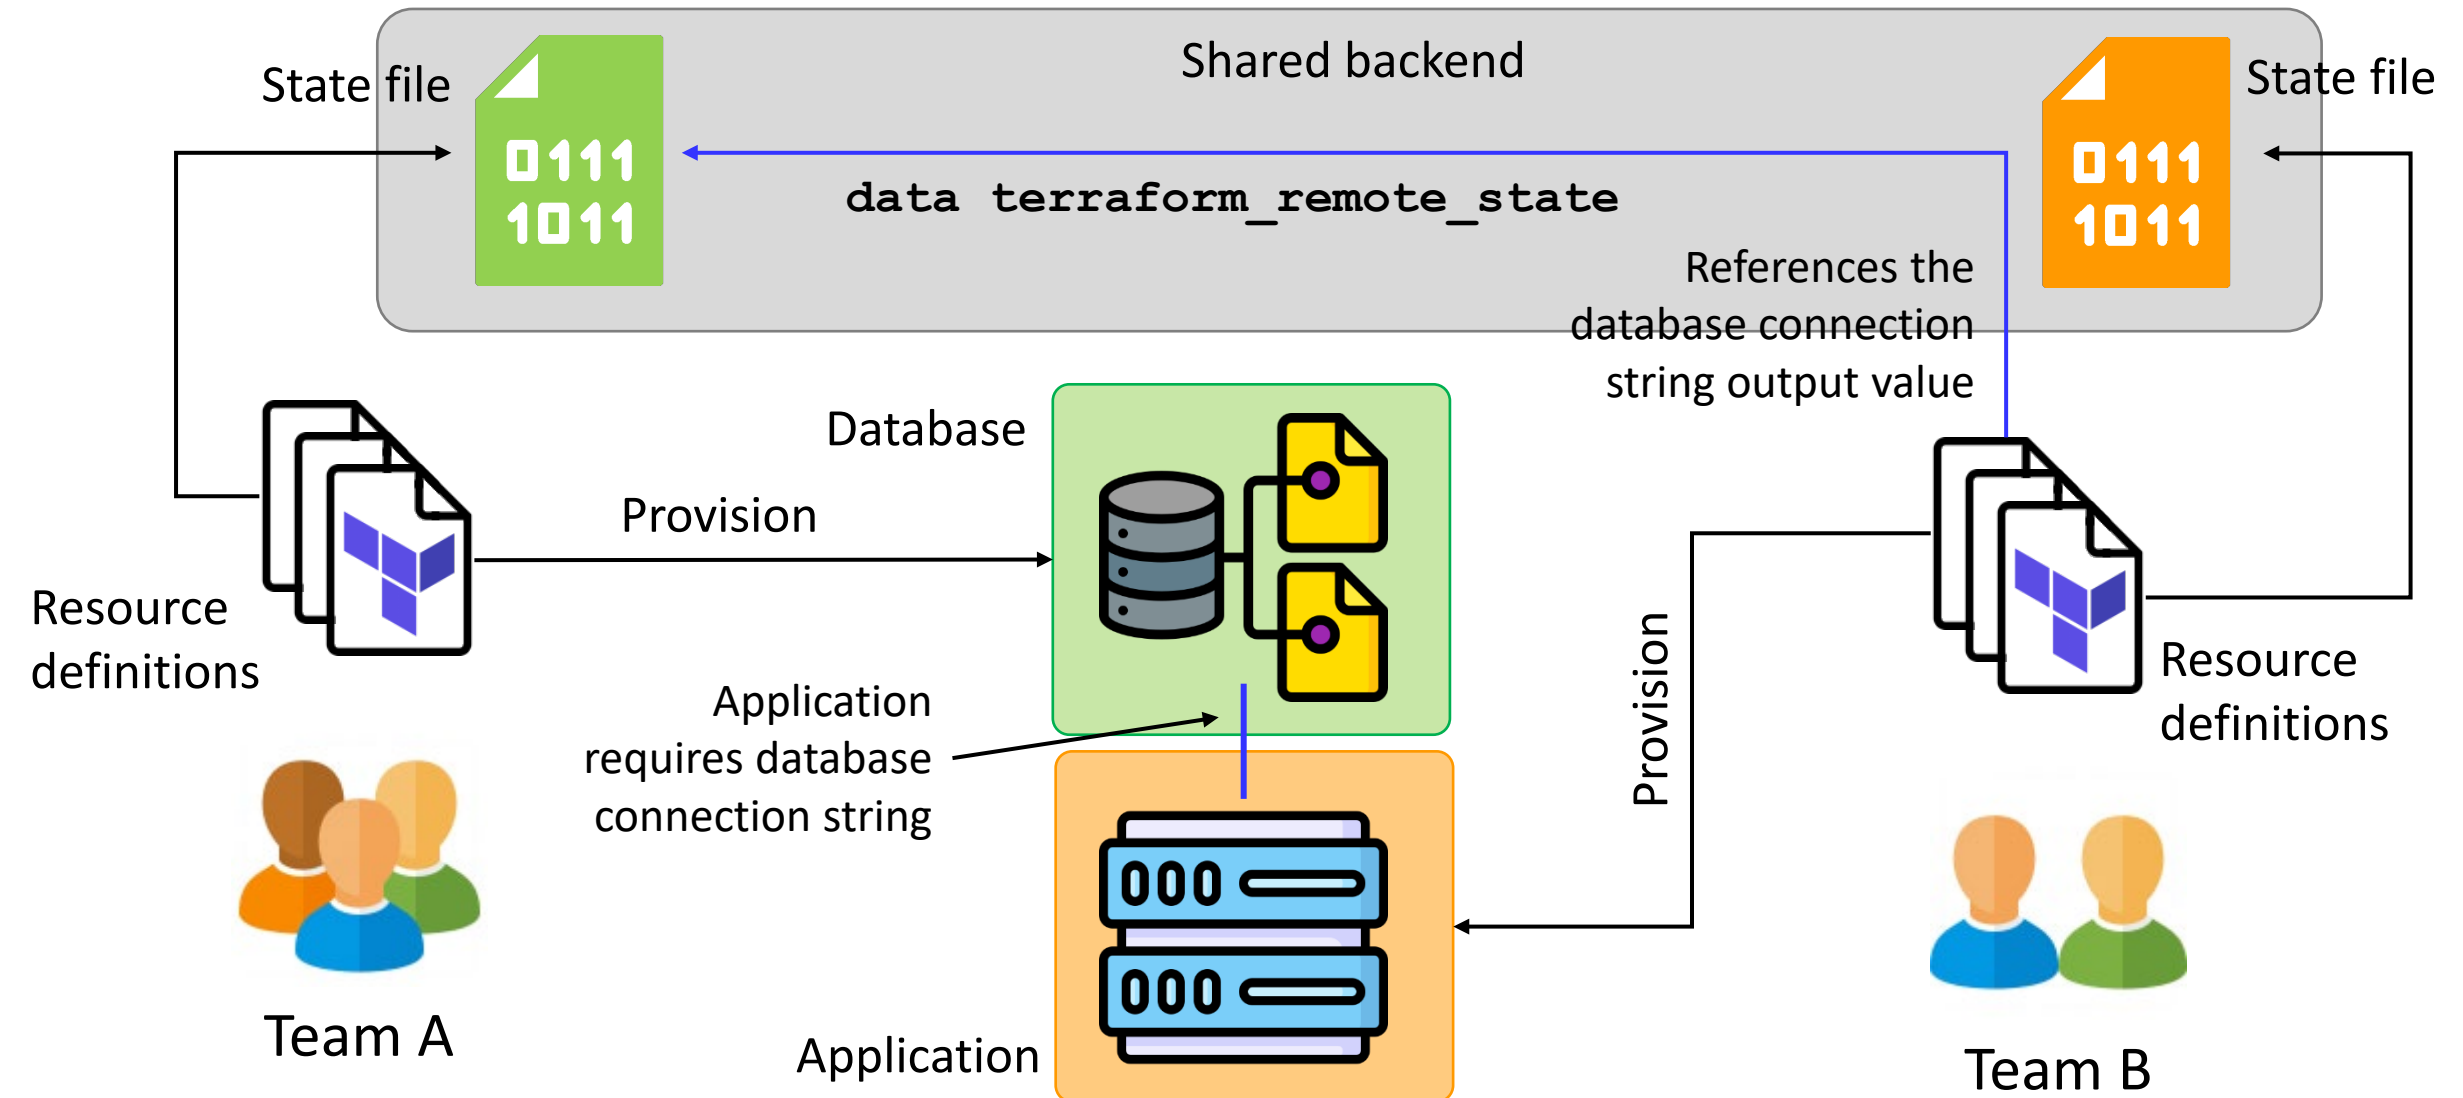


# Sharing State File

- Remotely accessible state files' output can be read by other Terraform scripts
  - Need to configure access for others to use
  - Can only access root output values
- Accessed with the `terraform_remote_state` data source
  - See <https://www.terraform.io/docs/language/state/remote-state-data.html>
- Scripts using the data source will not be notified if the data source is deleted



# Sharing and Accessing State File





# Example - Sharing State File

```
data terraform_remote_state db {  
  backend = "s3"  
  config = {  
    ...  
  }  
}
```

Type of backend

Configure the remote data source.  
Configuration should be the same as the  
backend

```
resource digitalocean_app_web {  
  spec {  
    name = "web"  
    region = var.region  
    env {  
      key = "DB_URL"  
      value = data.terraform_remote_state.db.outputs.db_connection  
      scope = "RUN_TIME"  
      type = "GENERAL"  
    }  
    service { ... }  
  }  
}
```

Can only access root project outputs viz. not  
module outputs used within the project  
db\_connection must be an output

The provisioned databased is  
used by another project



# No Variables in Backend Block

- Backend blocks cannot reference variables
  - See <https://www.terraform.io/docs/language/settings/backends/configuration.html#using-a-backend-block>
- Alternative
  - Use environment variables
    - Eg S3 - `AWS_REGION`, but not for all arguments
  - Use a variable file to pass to Terraform
    - Partial configuration



# Example - Partial Configuration

```
providers.tf
terraform {
  required_version = ">= 0.15.0"
  required_providers = {
    digitalocean = { ... }
  }

  backend "s3" {
    key = "terraform.tfstate"
  }
}
```

```
s3_backend.hcl
bucket = "mybucket"
region = "sgp1"
access_key = "...."
secret_key = "...."
```

All arguments required  
by specific backend

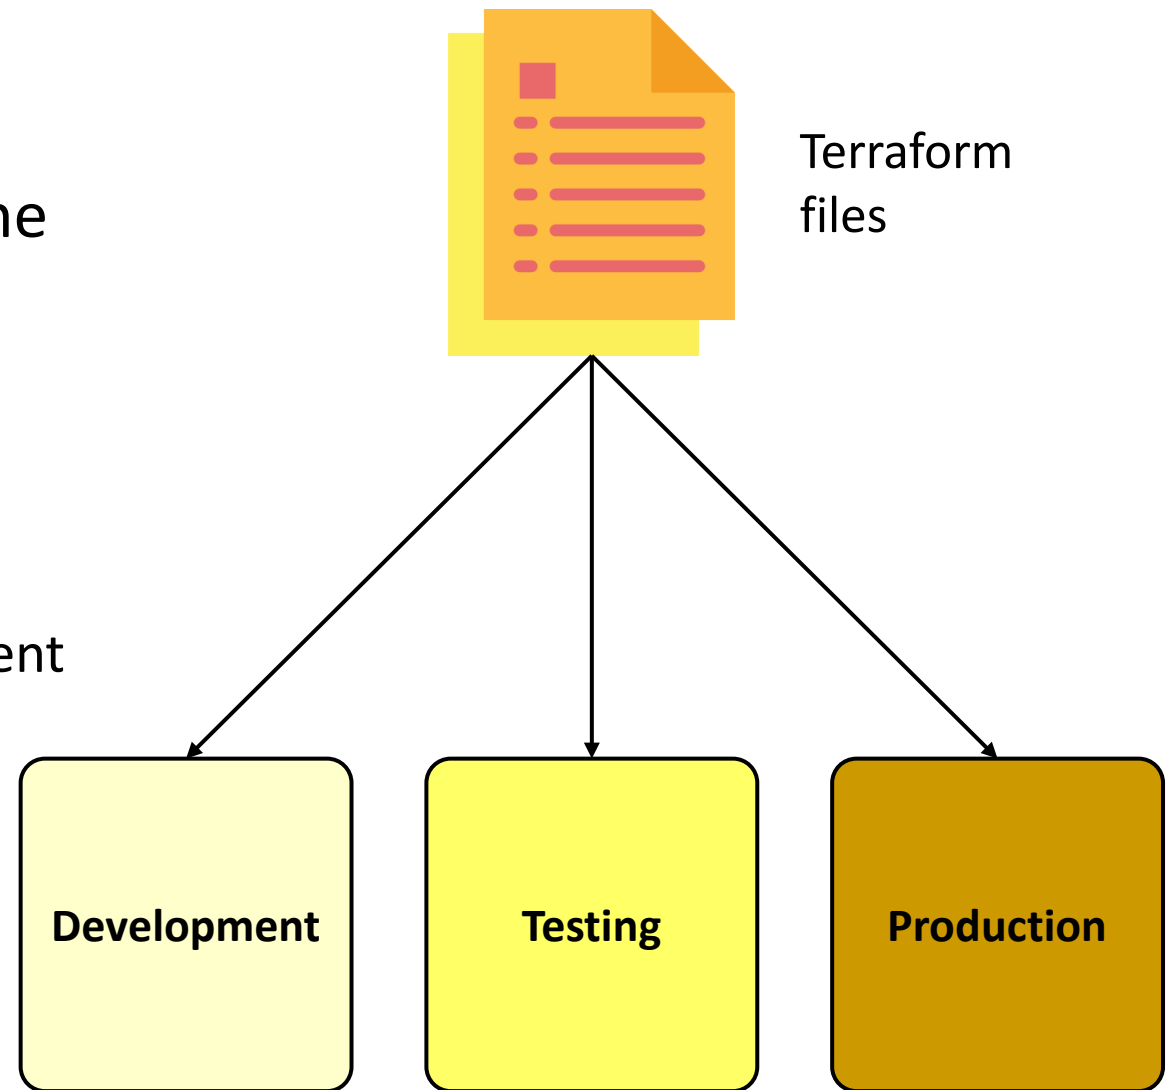
```
terraform init -backend-config=s3_backend.hcl
```

Initialize Terraform (with backend) by providing .hcl file



# Isolating Environment

- Sharing state file allows every one in the team to use the same environment
  - Share provisioned resources
- Separate environment so one environment does not interfere with another
  - Eg. testing a new OS version in development causes the production instances to be re-provisioned with the new version
- Two ways to isolate environment
  - Workspace
  - Directory structure





# Isolating Environment with Workspaces

- Workspaces allow you to create different state files in different workspaces
  - Terraform will report that it will have to provision all resources in a newly created workspace
- Terraform CLI will read all the `.tf`, `.tfvars` files from the current directory
  - Unless you edit the files before running apply, you will get the same set of resources
  - Not way to pass to the CLI a subset of the files from your current directory
- Error prone because if you forget to switch workspace before performing apply, can destroy the existing provisioned resource from the current workspace
- Workspace name as variable  
`terraform.workspace`

- List workspace

```
terraform workspace list
```

- Create a new workspace and switch to it

```
terraform workspace new myspace
```

- Change to an existing workspace

```
terraform workspace select otherspace
```

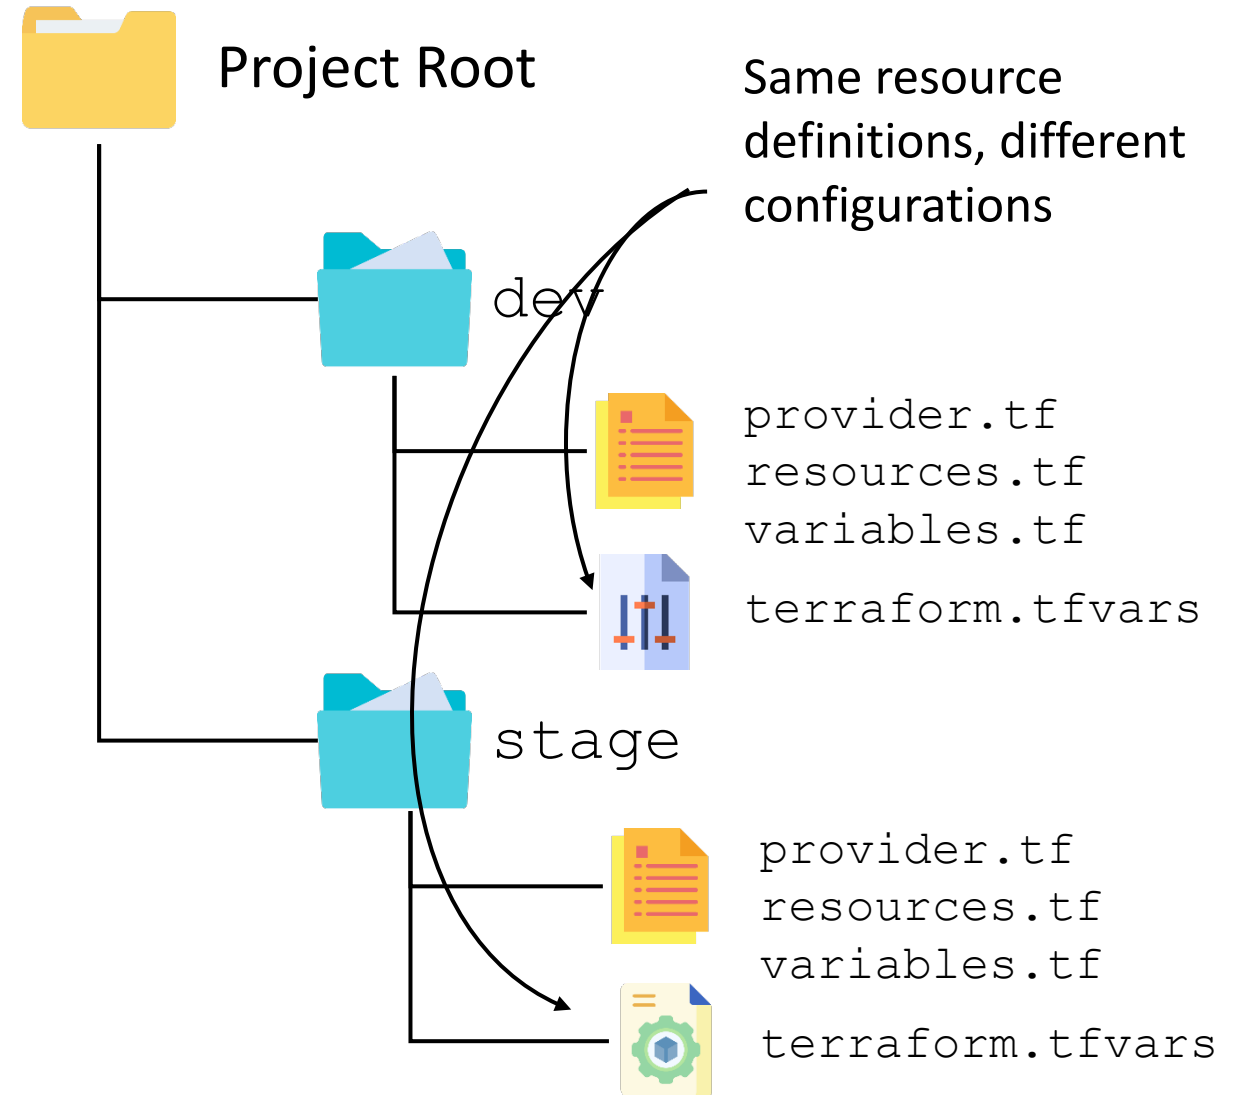
- Delete a workspace

```
terraform workspace destroy otherspace
```



# Isolating Environment with Directory Structure

- Create a directory for each environment eg. dev, test, prod
- Each directory will contain
  - Provider and resource definition files
  - Variable files specific to each environment
- Provider and resource definition are typically common across all the environment
- Better approach than workspaces but lots of duplication
- Reduce duplication by modularizing common and complex components

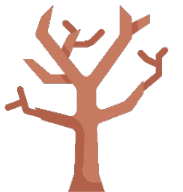






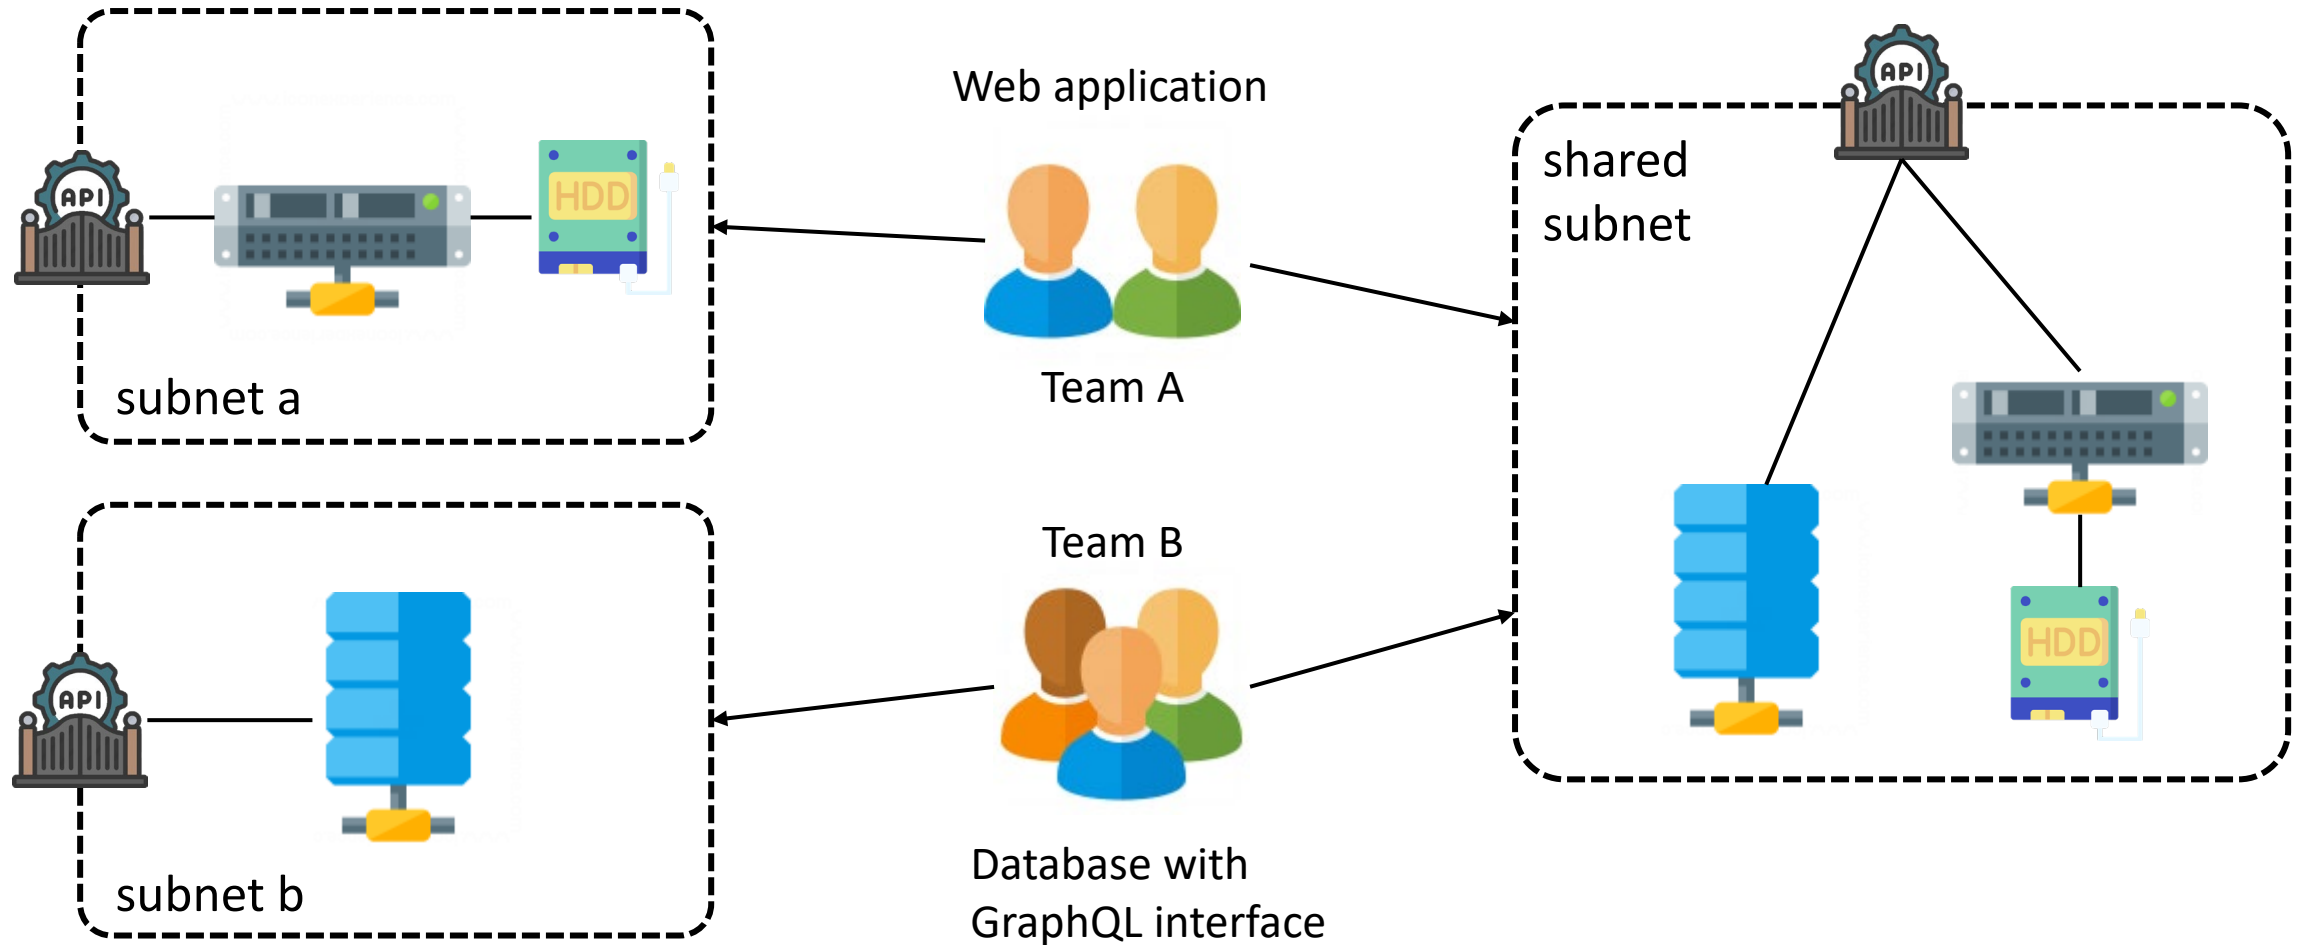
# Shared Elements - Resource Definitions

- Configuration, resources and infrastructure definition
  - Eg database connection configs, VM details, layout of the infrastructure/stacks
  - Slightly different in each team
- Handle by duplicating the files or parameterizing the shared elements
- Duplicating avoids tight coupling
  - WET - Write Everything Twice
  - Resource definitions are simpler - do not need to handle multiple scenerios
  - May diverge from prescribed company policy for using and consuming cloud resources
    - Eg. only use company's encrypted images
- Parameterizing avoids duplication
  - DRY - Don't Repeat Yourself
  - Treat resource definitions like APIs
  - Best practices can be shared across the teams





# Shared Elements - Environment





# Shared Elements - Environments

- Teams may have different environments for development/testing/deployment
  - No sharing, clean, do not interfere with other teams
  - Can be expensive, duplicate resources
- Provisioning into the same environment
  - Leverage share resources like internet gateways, bastion host, etc.
  - Consistent eg. security setting, routing rules, etc
- Sharing stateful resources
  - Servers - provisioning a new server or use an existing instance
    - Eg for deploying application - Docker or Kubernetes cluster



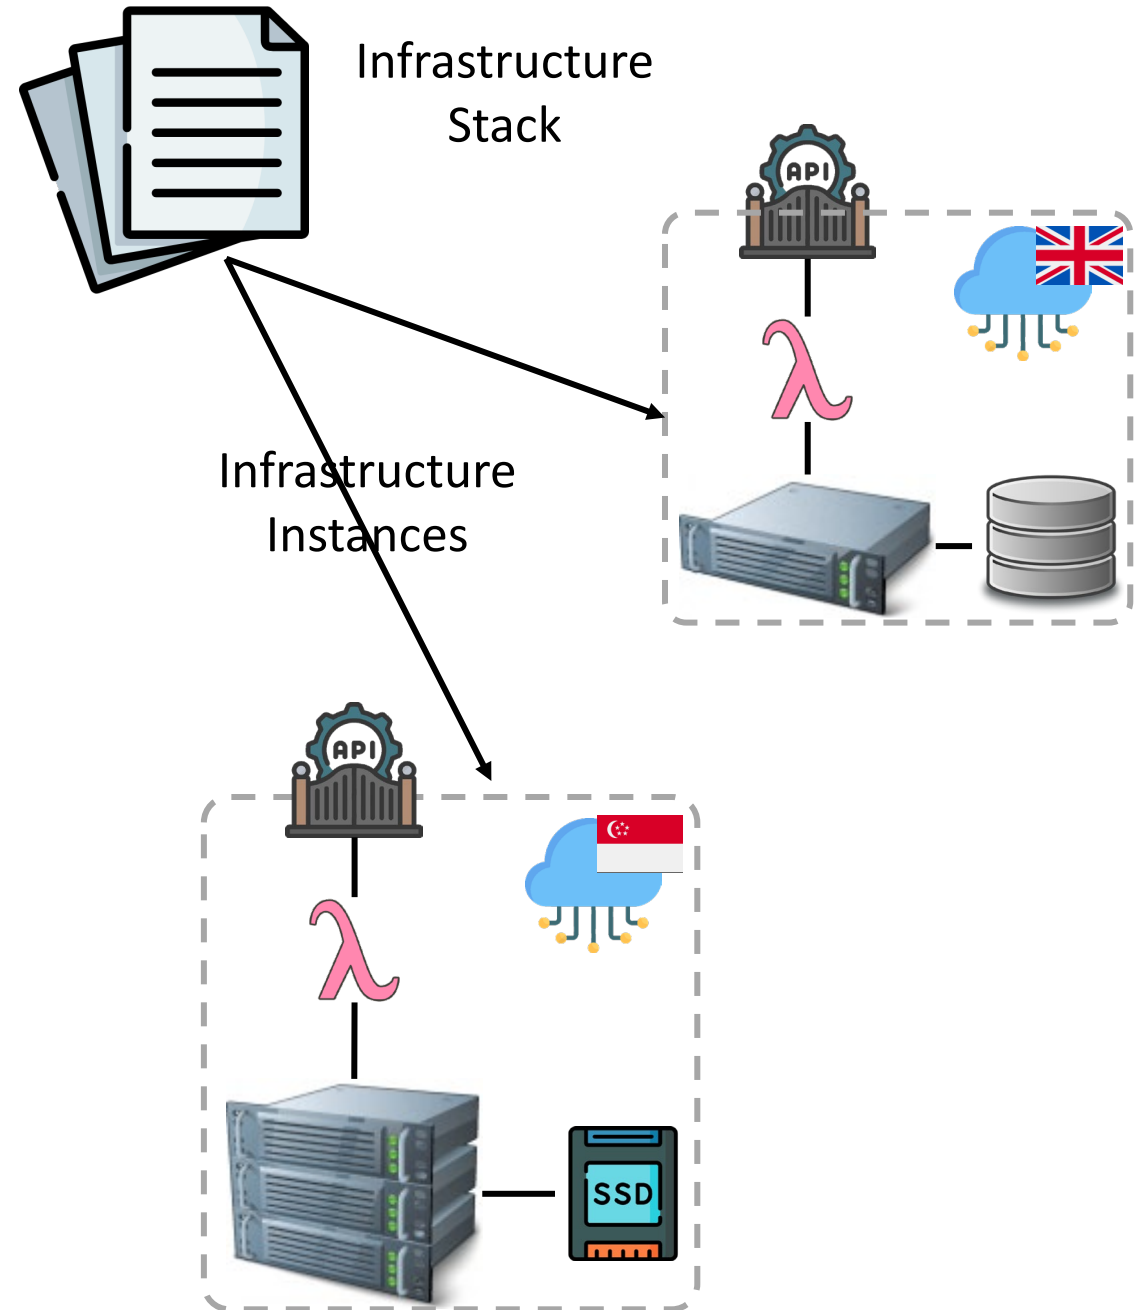
# Shared Elements - Environments

- Sharing stateful resources
  - Within a network, partition into subnets
    - Eg. Network - 10.10.10.0/24
    - Eg. Subnet - 10.10.10.0/26 - 4 subtnets, 62 host per subnet



# Infrastructure Stack

- Infrastructure stack includes
  - Images created from Packer
  - Provisioning with Terraform
  - Configured using Ansible
  - Templates for Terraform and Ansible
- Independent and parameterized to create the same infrastructure with different configuration
- Produces artefacts that can be used
  - Directly by the 'root' stack
  - As input to other stacks



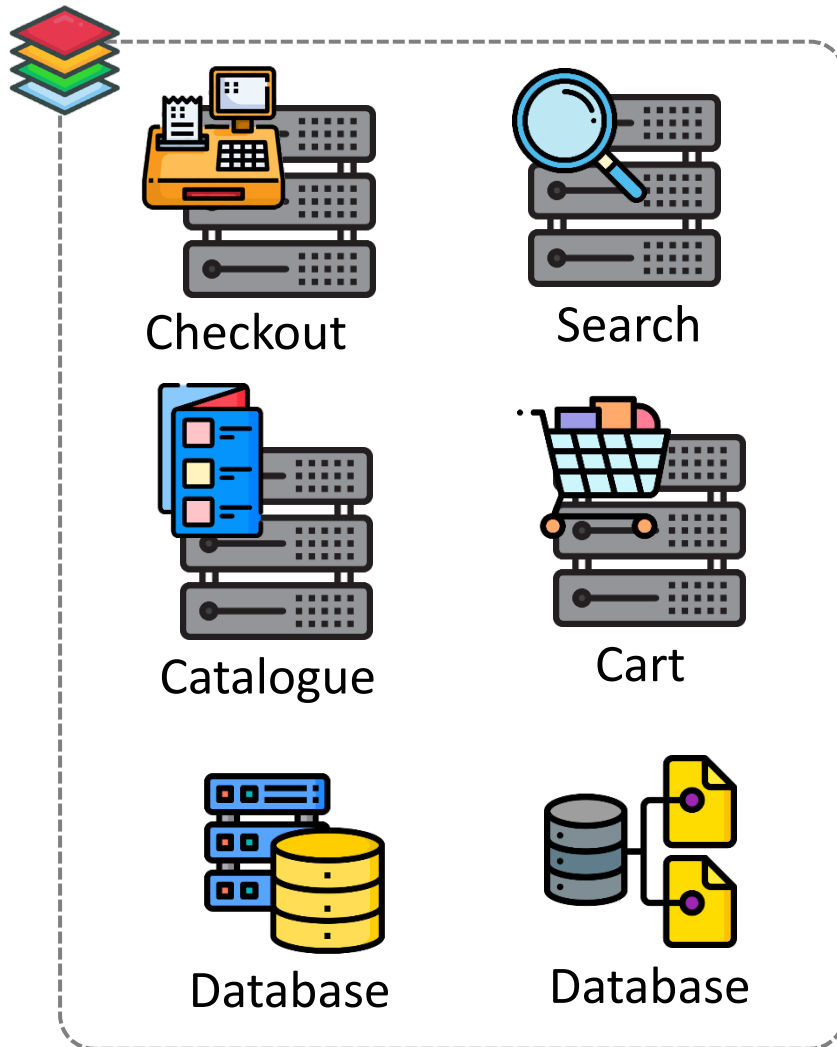


# Architectural Patterns

- Architectural patterns is common design used for infrastructure
  - Includes software installed and configurations
  - Similar in idea to design patterns for software design
- Infrastructure architectural patterns
  - Monolithic
  - Application group
  - Service
  - Micro



# Monolithic Stack

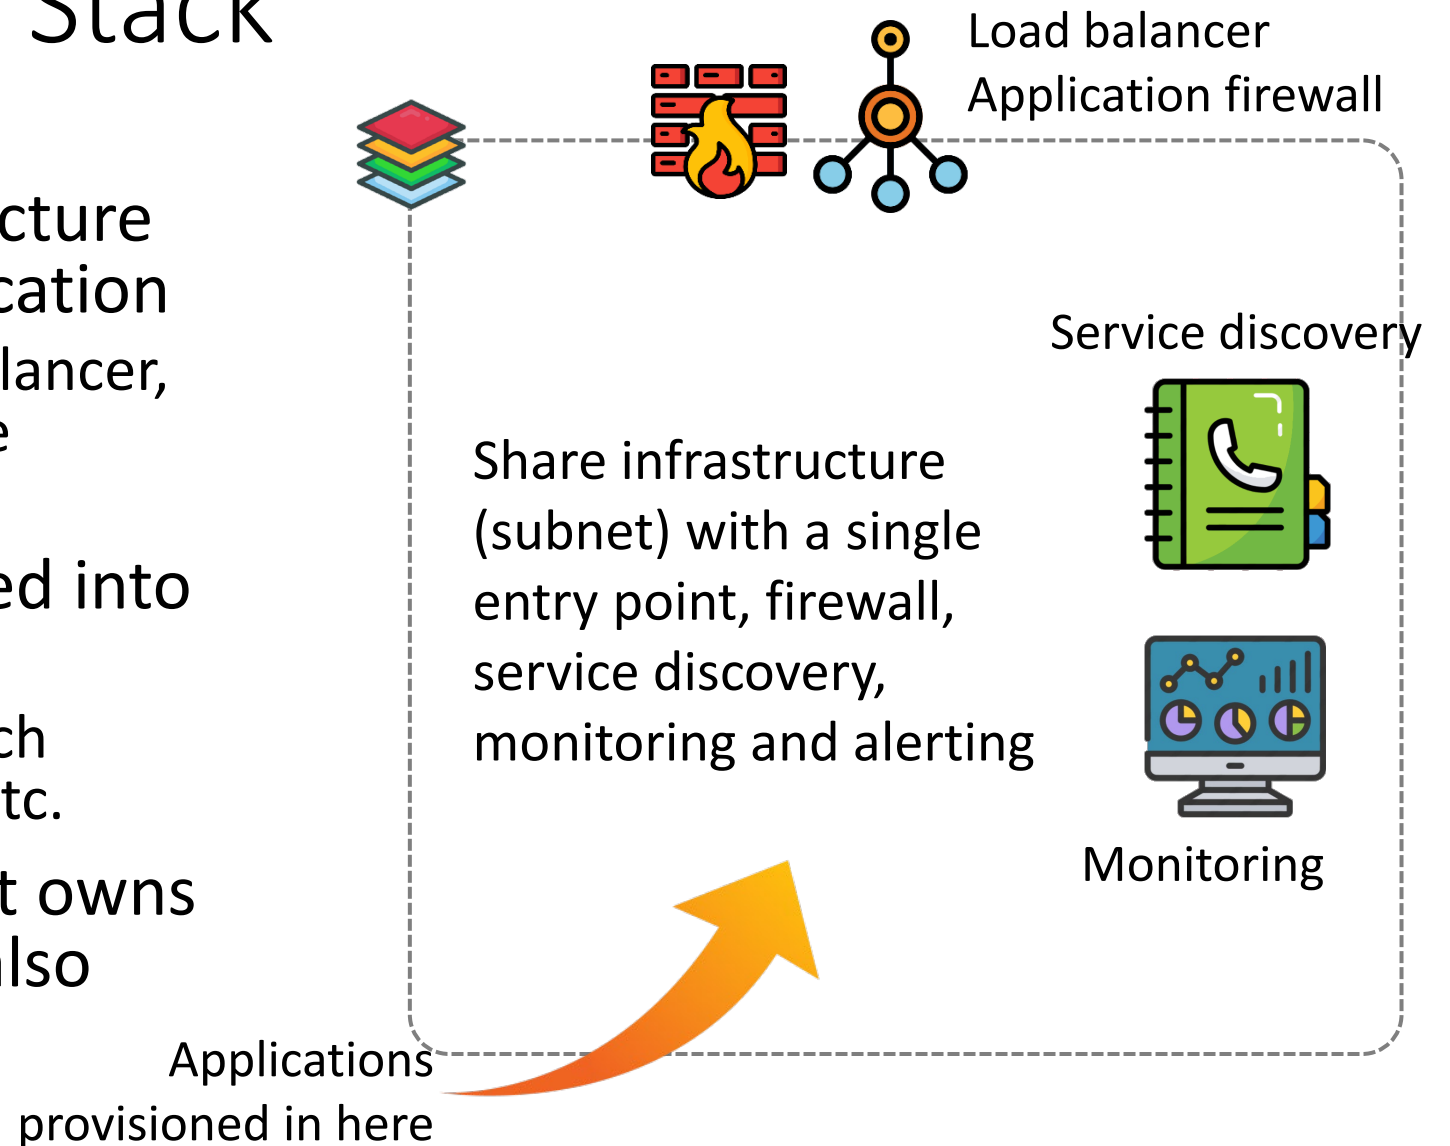


- All in one, easy to start
  - Easy to start, when grow organically will contain too many infrastructure elements
- When the stack grows
  - Difficult to separate the relationship within the
  - May take longer time to validate and provision
  - Can cause widespread changes with `terraform apply` when a change is made to an element
- Key indicator that a stack is monolithic is the manual coordination between members in the team



# Application Group Stack

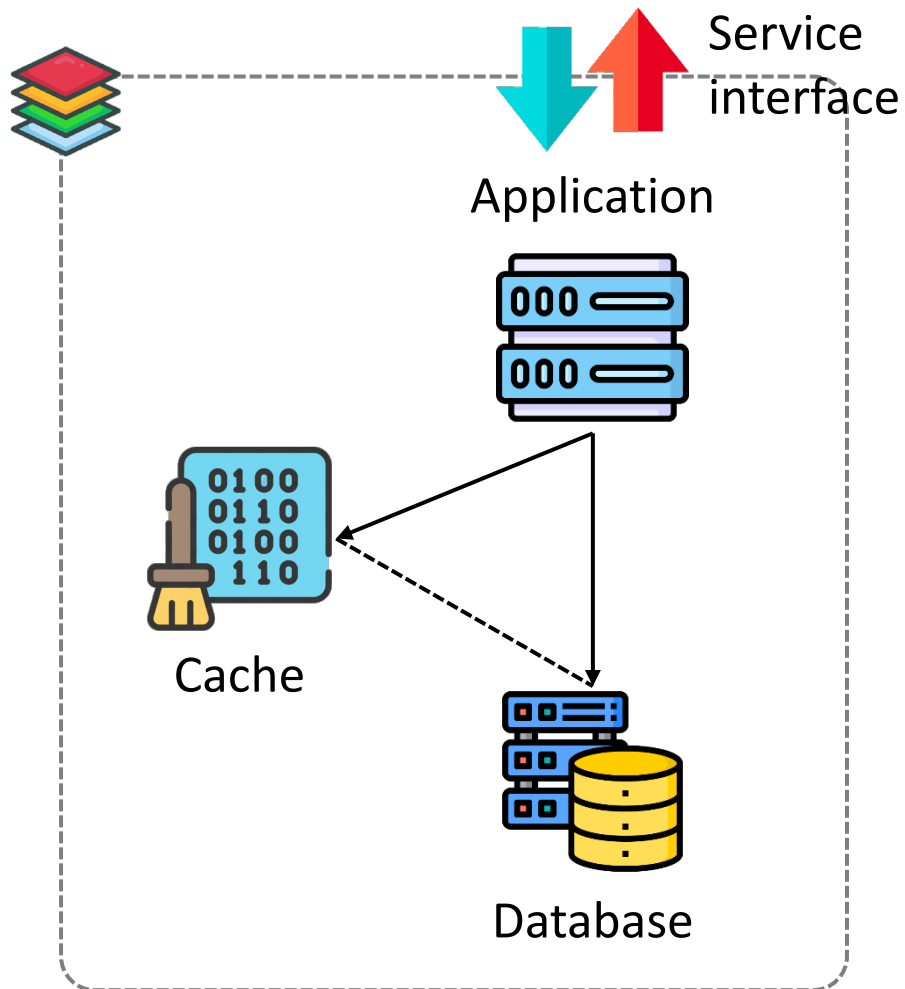
- Create a stack for infrastructure shared by a group of application
  - Resources like VPC, load balancer, application firewall, etc. are common and shared
- Applications are provisioned into the shared infrastructure
  - Eg. catalogue service , search service, checkout service, etc.
- Works well if the team that owns the shared infrastructure also owns the application







# Service Stack

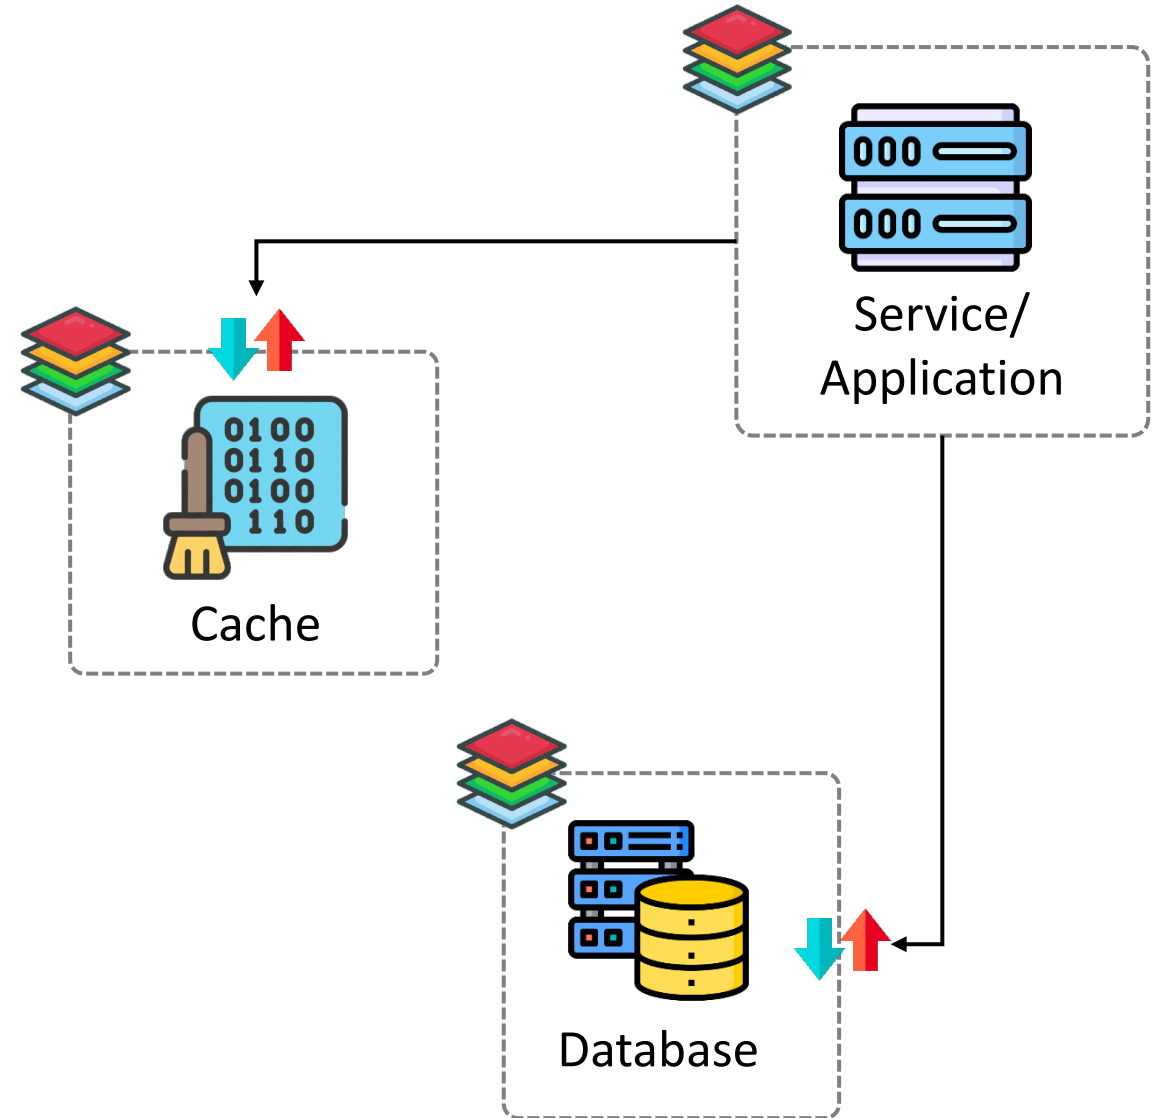


- Create a stack for each service
- Service stack may include
  - Compute (eg VM) for running the service/application
    - Service can be installed with Ansible or 'burnt' directly onto the image
  - Database to be used exclusively by the service
- Well define interface to access the service
- Works well with application group stack



# Micro Stack

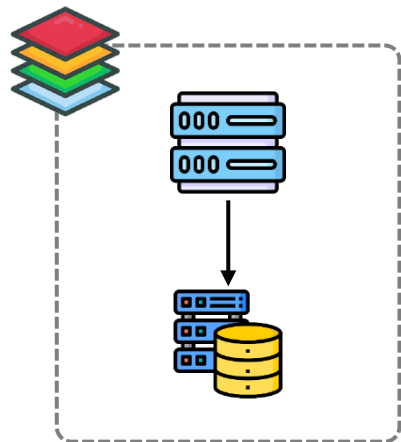
- Divides a single service across multiple stacks
- More moving parts, more complex
- Better availability when service is important or complex
  - Each stack can be configured and managed separately by expertise in each of the domain
  - Allow for different lifecycles for each of the service's components



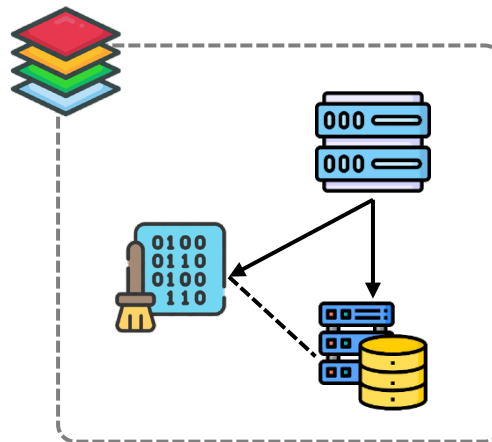


# Environments

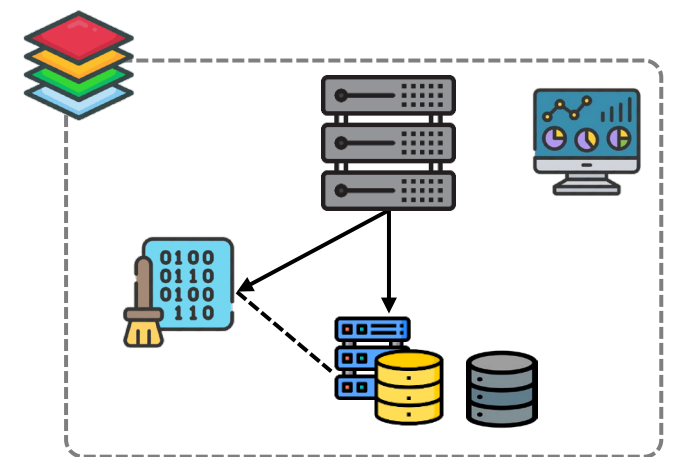
- Environments are resources and software used for a particular purpose
  - Eg. testing, development, staging, production
- Stack elements remains the same but the configuration of these elements may be slightly different
  - Eg. server size, clustered vs single instance databases, etc.
  - Eg. additional elements that may not be required eg. backups



Development



Testing



Production



# Appendix



# Reuse Common Task

- Keep common tasks in their separate file
  - Eg. upgrade all packages to the latest version
- Task file can be included in plays with `include_tasks`
- No explicit mechanism to pass parameters into the task
  - Task files can access all the variables defined in a play
- `import_tasks` is similar to `include_tasks` but
  - `import_tasks` is static - the tasks in the give YAML file is read and parsed when the playbook is parsed
  - `include_tasks` is dynamic - the tasks in the given YAML file is read and parsed when the playbook is executed
    - Allow different tasks to be include dynamically



# Example - import\_tasks

```
tasks:
- name: Update and install new packages
  import_tasks: tasks/apt_update_install.yaml
  var:
    package_list:
      - mysql-server
      - wordpress
      - php
      - libapache2-mod-php
      - php-mysql
```

List of task to  
be added to a  
playbook


Variable from playbook can be  
accessed by the task file

```
tasks/apt_update_install.yaml
- name: Update cache
  apt:
    update_cache: yes
- name: Upgrade installed packages
  apt:
    upgrade: yes
- name: Install new packages
  apt:
    name: '{{ item }}'
    state: latest
  when: package_list is defined
  loop: '{{ package_list }}'
```



# Example - include\_tasks

```
tasks:
- name: Update and install new packages
  include_tasks: 'tasks/{{ ansible_pkg_mgr }}.yaml'
var:
  package_list:
    - mysql-server
    - wordpress
    - php
    - libapache2-mod-php
    - php-mysql
```



The expression is dynamically evaluated when the playbook is executed  
Determine either to use `apt.yaml` or `yum.yaml` file to install the package depending on the package manager of the target host



# Roles

- Playbook is used to configure a host for a particular purpose
  - Eg. database server, bastion host, CMS/WordPress, etc.
  - Eg. Need to configure MySQL server multiple times in different projects
- A role is a set of task applied to a host to configure the host for a specific purpose
  - Provide an intuitive role name to identify the purpose of the tasks
  - Hide complex configuration by providing a way to reuse playbooks
  - Share know-how by reusing tasks
- It is independent and can be incorporated into any playbooks





# Using a Role

- Search from Ansible Galaxy, filter by roles
  - <https://galaxy.ansible.com/>
- Community, can publish the roles you have created
  - Require a GitHub account
- Documentation on how to use the role in your playbook
  - Version and OS platform
- Roles are executed before tasks
- Instructions to download and install a role
  - `ansible-galaxy`

```
ansible-galaxy install geerlingguy.docker
```

Install Docker role  
from Ansible Galaxy

<https://galaxy.ansible.com/geerlingguy/docker>



# Example - Using the Docker

site.yaml

```
- name: Install Docker
  host: docker_host
  pre_tasks:
    - name: Create user
      user:
        name: '{{ docker_user }}'
        password: '{{ docker_user_password }}'
  roles:
    - role: roles/docker
      vars:
        docker_users:
          - '{{ docker_user }}'
  tasks:
    - name: Enable remote access
      lineinfile:
        path: /lib/systemd/system/docker.service
        regexp: '^ExecStart='
        line: ExecStart=/usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock-
H=tcptcp://0.0.0.0:2375
    - name: Restart Docker daemon
      systemd:
        name: docker
        daemon_reload: true
        state: restarted
  ...
```

Like task but  
executed before  
roles or tasks

Role is downloaded  
to roles directory

Role variable(s) to  
configure Docker

Note: an alternative way to  
provision Docker host is to use  
docker-machine  
<https://docs.docker.com/machine>



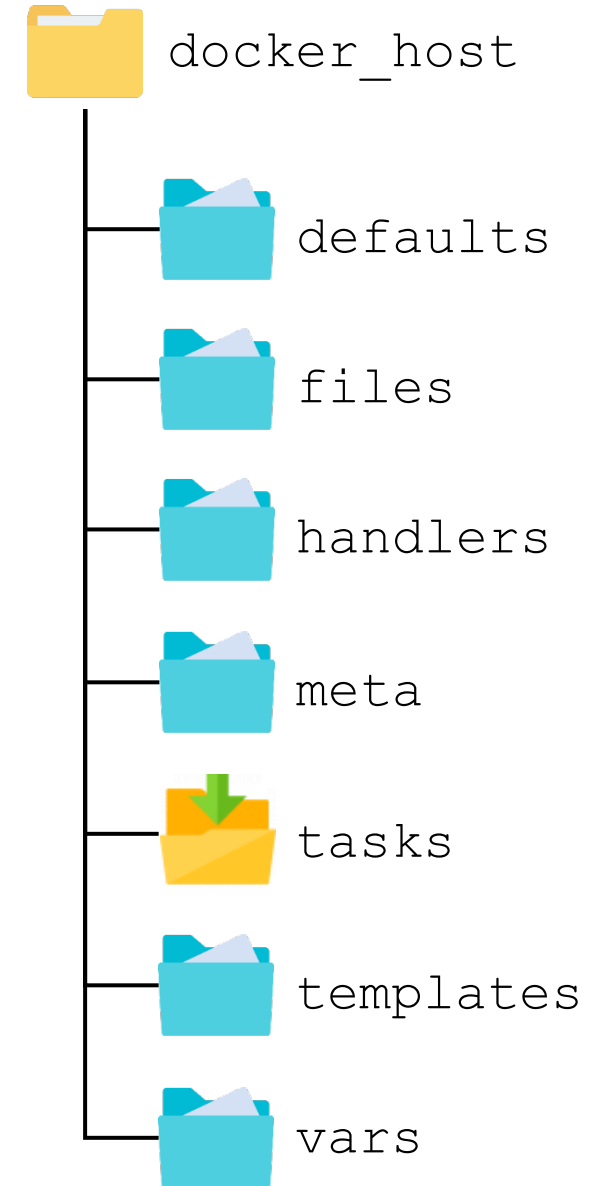
# Creating Role

- A set of pre defined directory that can be generated with ansible-galaxy

```
ansible-galaxy init docker_host
```

- **Directory uses**

- `defaults` - default variables
- `files` - files, eg configuration, to be copied over to the target
- `handlers` - for handlers
- `tasks` - contains all the task to be performed by the role. When a role is applied, Ansible will perform the list of task listed here
- `templates` - dynamically generated files
- `vars` - other variables, override defaults
- `meta` - metadata for the role





# Example - Docker Host Role

```
default/main.yaml
```

```
---
```

```
package_list:
```

- apt-transport
- ca-certificates
- software-properties-common
- curl
- openssh-server

```
docker_gpg_key: https://download.docker.com/linux/ubuntu/gpg
```

```
docker_repo: "deb [arch=amd64] https://download.docker.com/linux/ubuntu  
focal stable"
```

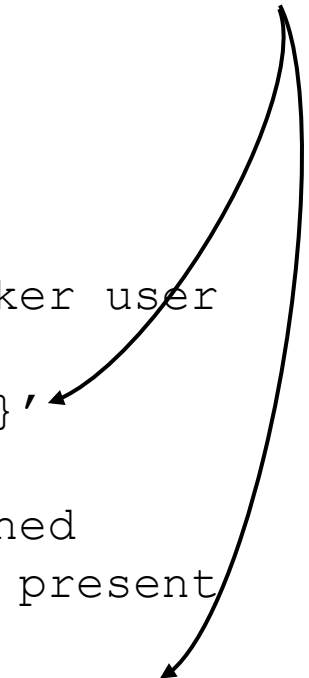


# Example - Docker Host Role

```
tasks/main.yaml
---
- name: Install packages
  apt:
    update_cache: yes
    name: '{{ item }}'
    state: latest
  loop: '{{ package_list }}'
- name: Add Docker GPG key
  apt-key:
    url: '{{ docker_gpg_key }}'
    state: present
- name: Add Docker repository
  apt_repository:
    repo: '{{ docker_repo }}'
    state: present
```

User defined variables

```
- name: Install Docker
  apt:
    update_cache: yes
    name: docker-ce
    state: latest
- name: Create non-root Docker user
  user:
    name: '{{ docker_user }}'
    groups: [ 'docker' ]
    when: docker_user is defined
- name: Check if SSH key is present
  stat:
    path: '{{ docker_ssh_key }}'
    register: docker_ssh_key is defined
- name: Add SSH key to user
  authorized_key:
    user: '{{ docker_user }}'
    key: '{{ lookup('file', docker_ssh_key )}}'
    when: ssh_key.stat.exists
```





# Example - Using Docker Host Role

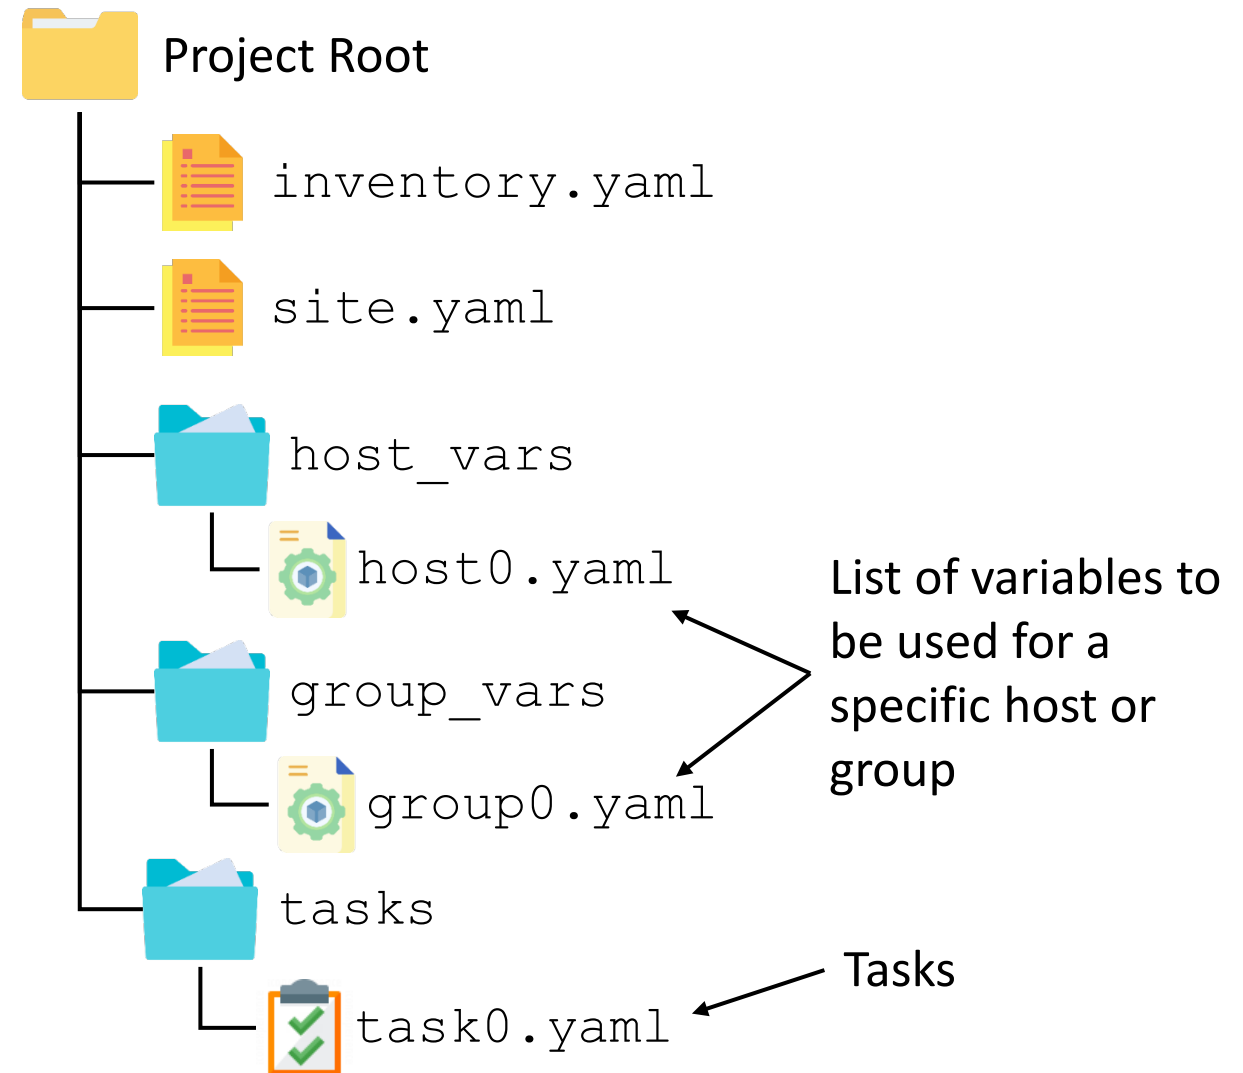
site.yaml

```
- name: Install Docker
  hosts: a_docker_host
  roles:
  - role: roles/docker_host
  vars:
    docker_user: fred
    docker_ssh_key: /path/to/public/key
```



# Structuring an Ansible Project

- Single playbook cannot scale
- Use separate specially named directories for hosts and groups
- Put commonly used task in a separate file so that playbooks can reuse them





# host\_vars and group\_vars Directories

- Special directories to define host and group variables
- Keep the number of variable declarations in inventory file to a minimum
- File name is the name of the host or group
  - Eg. `hosts_vars/localhost.yaml`, `group_vars/all.yaml` for variables to be applied to localhost and the all group
- When a playbook is executed, the variables from these files will be merged with those from the inventory file (if any)





# Example - host\_vars and group\_vars

```
all:
  hosts:
    server-0:
      ansible_host: 192.168.0.100
      ansible_connection: ssh
      ansible_user: fred
      ansible_password: fred
    server-1:
      ansible_host: 192.168.0.101
      ansible_connection: ssh
      ansible_user: fred
      ansible_password: fred
    server-2:
      ansible_host: 192.168.0.102
      ansible_connection: ssh
      ansible_user: fred
      ansible_password: fred
      db_user: barney
      db_password: barney
      db_name: inventory
```



# Example - host\_vars and group\_vars

**inventory.yaml**

**all:**

hosts:

**server-0:** {}

**server-1:** {}

**server-2:** {}

**group\_vars/all.html**

ansible\_connection: ssh

ansible\_user: fred

ansible\_password: fred

**host\_vars/server-0.yaml**

ansible\_host: 192.168.0.100

**hosts\_vars/server-1.yaml**

ansible\_host: 129.168.0.101

**hosts\_vars/server-2.yaml**

ansible\_host: 129.168.0.102

db\_user: barney

db\_password: barney

db\_name: inventory



# Roles

- Playbook is used to configure a host for a particular purpose
  - Eg. database server, bastion host, CMS/WordPress, etc.
  - Eg. Need to configure MySQL server multiple times in different projects
- A role is a set of task applied to a host to configure the host for a specific purpose
  - Provide an intuitive role name to identify the purpose of the tasks
  - Hide complex configuration by providing a way to reuse playbooks
  - Share know-how by reusing tasks
- It is independent and can be incorporated into any playbooks



# Using a Role

- Search from Ansible Galaxy, filter by roles
  - <https://galaxy.ansible.com/>
- Community, can publish the roles you have created
  - Require a GitHub account
- Documentation on how to use the role in your playbook
  - Version and OS platform
- Roles are executed before tasks
- Instructions to download and install a role
  - `ansible-galaxy`

```
ansible-galaxy install geerlingguy.docker
```

Install Docker role  
from Ansible Galaxy

<https://galaxy.ansible.com/geerlingguy/docker>



# Example - Using the Docker

site.yaml

```
- name: Install Docker
  host: docker_host
  pre_tasks:
    - name: Create user
      user:
        name: '{{ docker_user }}'
        password: '{{ docker_user_password }}'
  roles:
    - role: roles/docker
      vars:
        docker_users:
          - '{{ docker_user }}'
  tasks:
    - name: Enable remote access
      lineinfile:
        path: /lib/systemd/system/docker.service
        regexp: '^ExecStart='
        line: ExecStart=/usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock-
H=tcptcp://0.0.0.0:2375
    - name: Restart Docker daemon
      systemd:
        name: docker
        daemon_reload: true
        state: restarted
  ...
```

Like task but  
executed before  
roles or tasks

Role is downloaded  
to roles directory

Role variable(s) to  
configure Docker

Note: an alternative way to  
provision Docker host is to use  
docker-machine  
<https://docs.docker.com/machine>



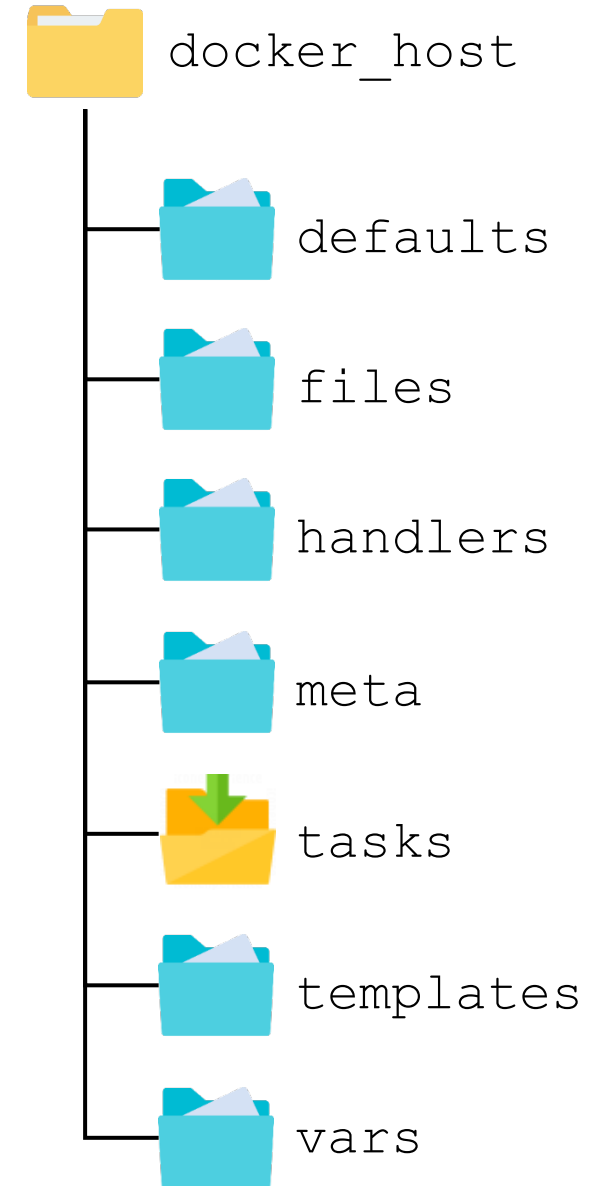
# Creating Role

- A set of pre defined directory that can be generated with ansible-galaxy

```
ansible-galaxy init docker_host
```

- **Directory uses**

- `defaults` - default variables
- `files` - files, eg configuration, to be copied over to the target
- `handlers` - for handlers
- `tasks` - contains all the task to be performed by the role. When a role is applied, Ansible will perform the list of task listed here
- `templates` - dynamically generated files
- `vars` - other variables, override defaults
- `meta` - metadata for the role





# Example - Docker Host Role

```
default/main.yaml
```

```
---
```

```
package_list:
```

- apt-transport
- ca-certificates
- software-properties-common
- curl
- openssh-server

```
docker_gpg_key: https://download.docker.com/linux/ubuntu/gpg
```

```
docker_repo: "deb [arch=amd64] https://download.docker.com/linux/ubuntu  
focal stable"
```

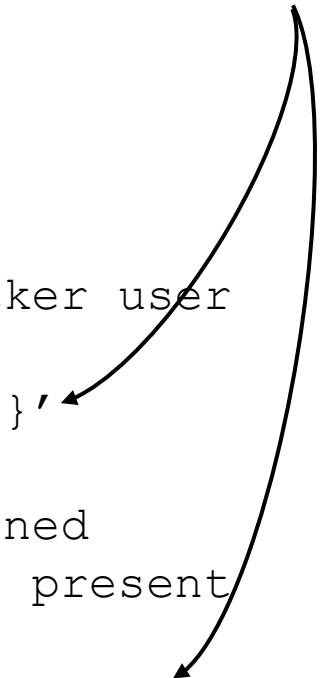


# Example - Docker Host Role

User defined variables

```
tasks/main.yaml
---
- name: Install packages
  apt:
    update_cache: yes
    name: '{{ item }}'
    state: latest
  loop: '{{ package_list }}'
- name: Add Docker GPG key
  apt-key:
    url: '{{ docker_gpg_key }}'
    state: present
- name: Add Docker repository
  apt_repository:
    repo: '{{ docker_repo }}'
    state: present
```

```
- name: Install Docker
  apt:
    update_cache: yes
    name: docker-ce
    state: latest
- name: Create non-root Docker user
  user:
    name: '{{ docker_user }}'
    groups: [ 'docker' ]
    when: docker_user is defined
- name: Check if SSH key is present
  stat:
    path: '{{ docker_ssh_key }}'
    register: docker_ssh_key is defined
- name: Add SSH key to user
  authorized_key:
    user: '{{ docker_user }}'
    key: '{{ lookup('file', docker_ssh_key ) }}'
    when: ssh_key.stat.exists
```







# Example - Using Docker Host Role

site.yaml

```
- name: Install Docker
  hosts: a_docker_host
  roles:
  - role: roles/docker_host
  vars:
    docker_user: fred
    docker_ssh_key: /path/to/public/key
```



# Example - HTTP/WebDAV

```
terraform {  
  required_version = ">= 0.15.0"  
  required_providers = {  
    digitalocean = { ... }  
  }  
}
```

The endpoint for creating, retrieving  
and updating the state file

Associate an identifier to  
endpoint if the server is  
used by many projects

```
backend http {  
  address = "http://myserver.com/myproj"  
  lock_address = "http://myserver.com/myproj"  
  unlock_address = "http://myserver.com/myproj"  
  username = "fred"  
  password = "fred"  
}  
}
```

Specify username and  
password if endpoint is uses  
HTTP Basic authentication

Locking is optional.  
Configure locking if  
server supports WebDAV



# Example - HTTP/WebDAV



Other members

```
LOCK /myproject
```

```
Authorization: Basic btoa("username:password")
```

```
<lock details in the body>
```

```
POST /myproject?lockid=...
```

```
Authorization: Basic btoa("username:password")
```

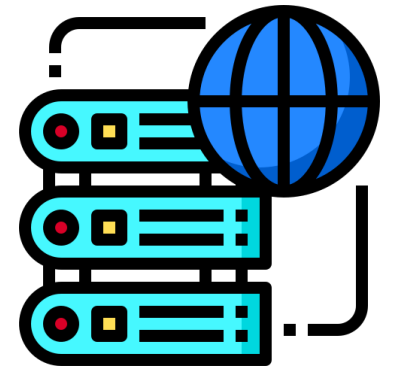
```
<update state file>
```

```
UNLOCK /myproject?lockid=...
```

```
Authorization: Basic btoa("username:password")
```



Developer



HTTP/WebDAV